



DeepLearning.AI

Why PyTorch?

Getting Started with PyTorch

Adding Two Numbers with PyTorch

```
result = a + b
```

```
a = torch.tensor(2.0)
b = torch.tensor(3.0)
result = a + b
print(result)
```

```
C:>
tensor(5.)
```

Machine
Learning



Traditional
Programming

Product Recommendation



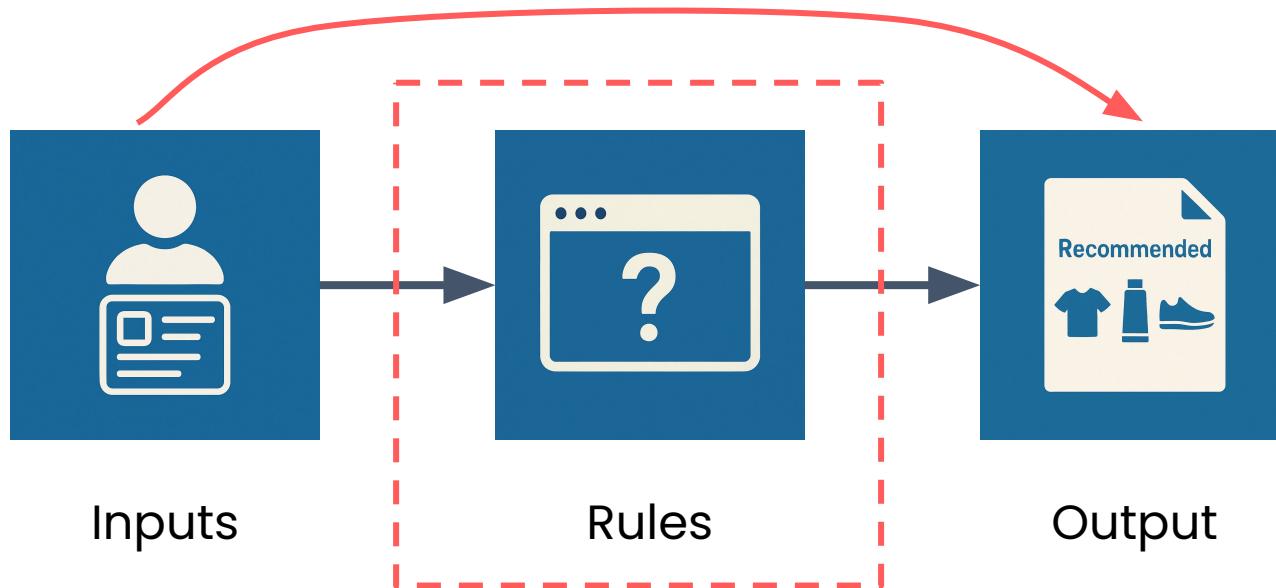
Product Recommendation



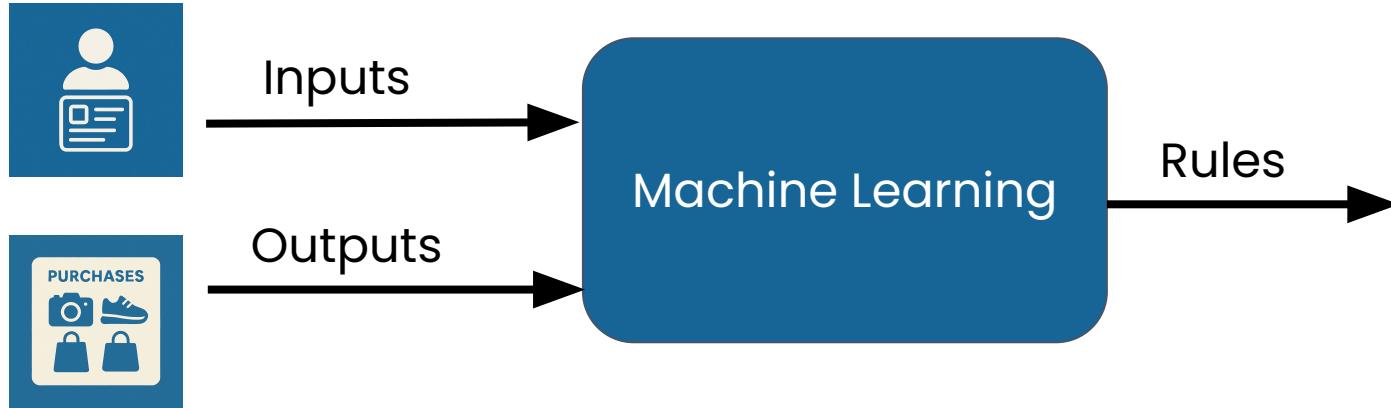
Rule: *"If a customer buys a camera, recommend camera lenses"*

- What if it was a gift?
- What if they already own five lenses?

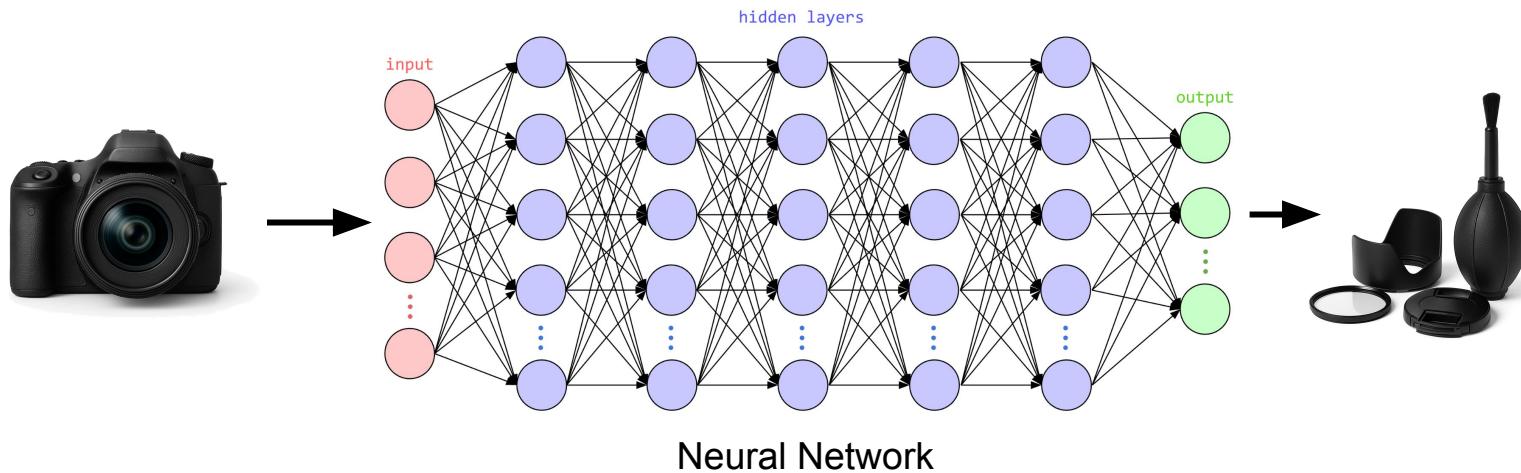
Traditional Programming



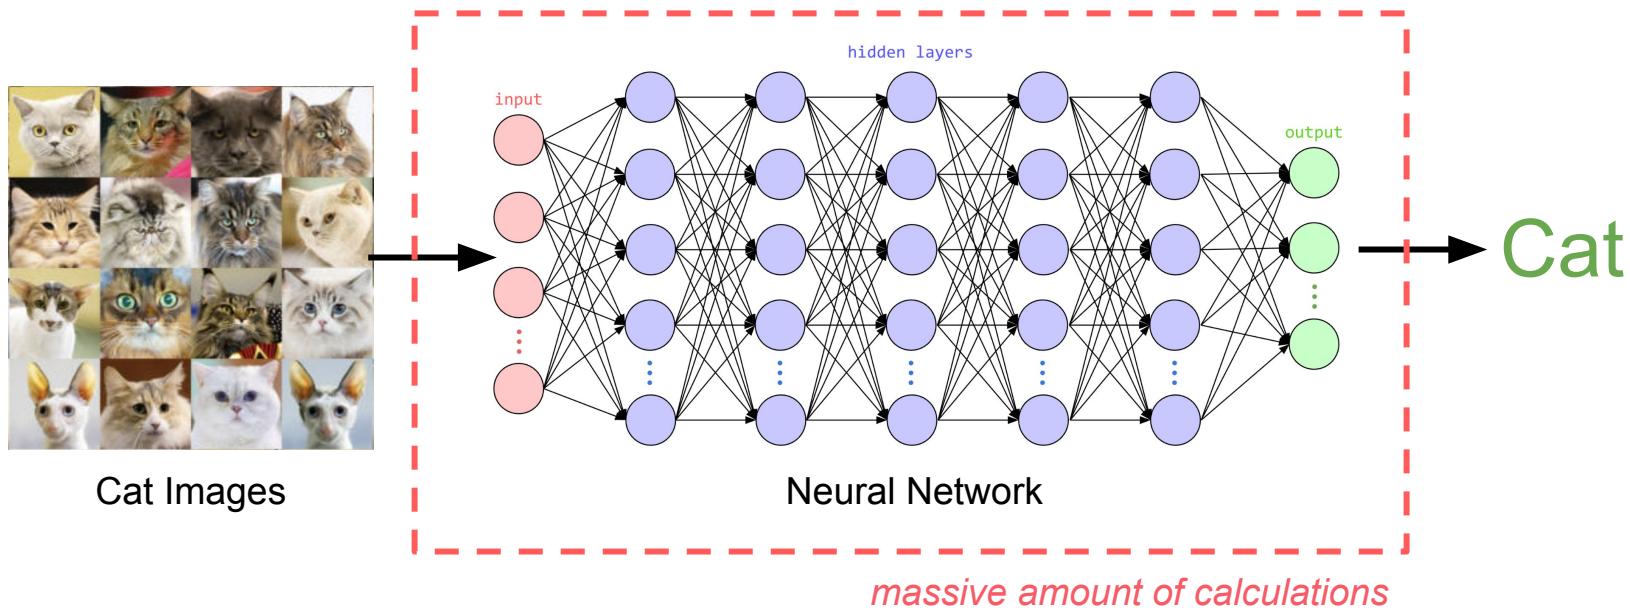
Machine Learning



Deep Learning – Neural Networks



Deep Learning – Neural Networks



Early Frameworks

```
result = a + b
```

```
a = framework.placeholder(float32)
b = framework.placeholder(float32)
result = framework.add(a,b)
with framework.Executor() as executor:
    actual_result = executor.run(result,
data_dictionary={a: 2.0, b: 3.0}))
```

```
C:>
<Object 'Add:0' shape=(2,) dtype=float32>
```

Early Frameworks

Static Computation Graph

Early Frameworks

```
# A chain of operations, one feeding into the next
result1 = do_math(input)
result2 = do_math(result1)
result3 = do_math(result2)
result4 = do_math(result3)
result5 = do_math(result4)
# ... and so on, down the line
```

Early Frameworks

```
# A chain of operations, one feeding into the next
result1 = do_math(input)
result2 = do_math(result1)

new_op = do_math(result2) ←

result3 = do_math(result2)
result4 = do_math(result3)
result5 = do_math(result4)
```

Early Frameworks

```
# A chain of operations, one feeding into the next
result1 = do_math(input)
result2 = do_math(result1)
new_op = do_math(result2)

result3 = do_math(result2)
result4 = do_math(result3)
result5 = do_math(result4)
```

No Testing of Parts

Early Frameworks

```
# A chain of operations, one feeding into the next
result1 = do_math(input)
result2 = do_math(result1)

new_op = do_math(result2)

result3 = do_math(result2)
result4 = do_math(result3)
result5 = do_math(result4)
```

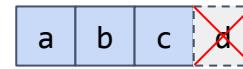
Full Build Required

Early Frameworks

~~if a:~~

~~for a in b:~~

Fixed input size



~~debugging~~

untraceable
system errors ~~X~~

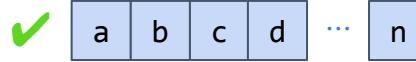


PyTorch

Early Frameworks

✓ If a:

✓ for a in b:



✓ debugging

✓ In code errors

✓ Python-like code

PyTorch Ecosystem



OpenMined

SKORCH



adver
torch

PyTorch



PyTorch
geometric



PENNYLANE

DeepLearning.AI

BoTorch

flair

Laurence Moroney



DeepLearning.AI

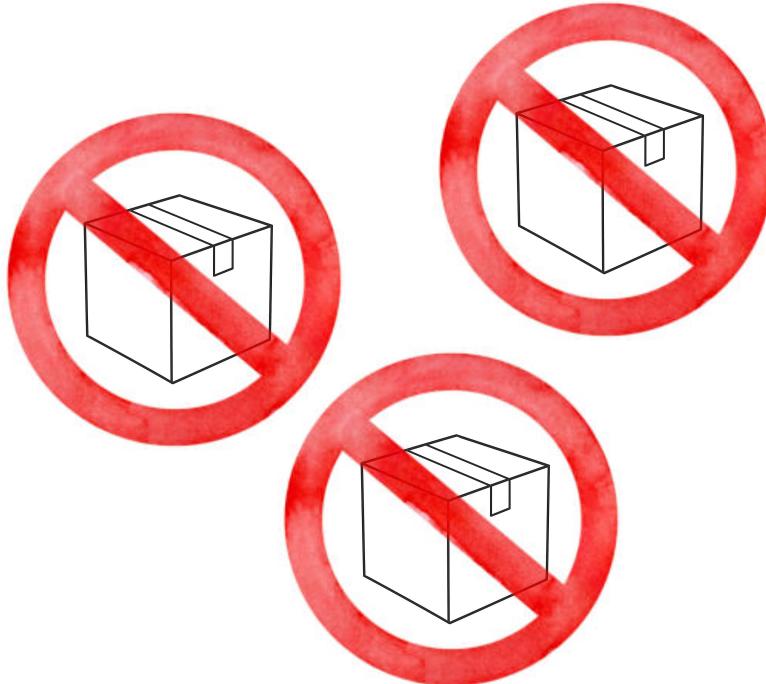
The Building Block of Neural Networks

Getting Started with PyTorch

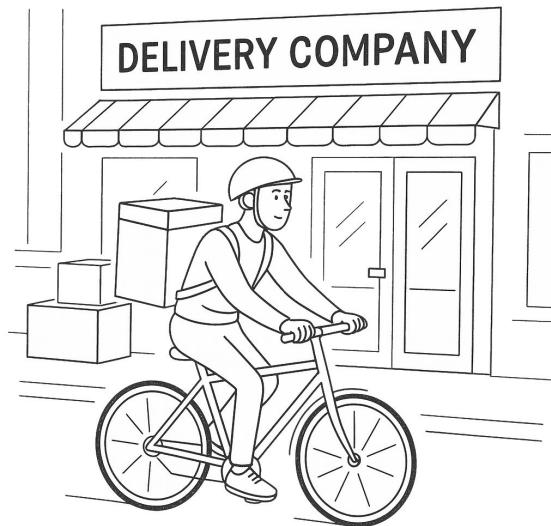
The Delivery Problem



30 minute guarantee

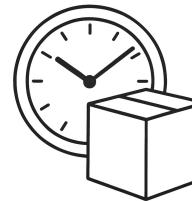


The Delivery Problem



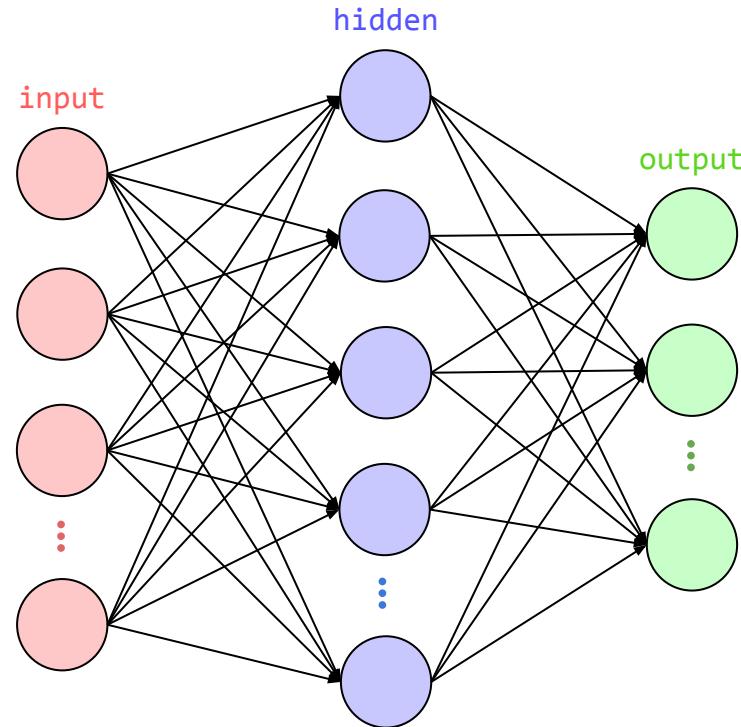
30 minute guarantee

Take the job?

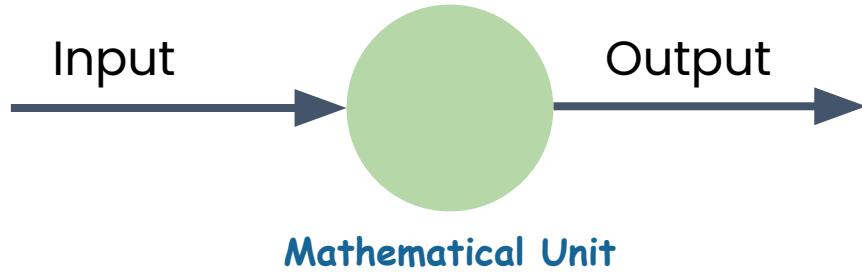


7 mile delivery 

Neural Networks



Single Neuron



Delivery Data

Distance (miles)	Delivery Time (minutes)
2.0	11.9
3.0	15.3
4.0	18.8
5.0	22.2
6.0	25.6

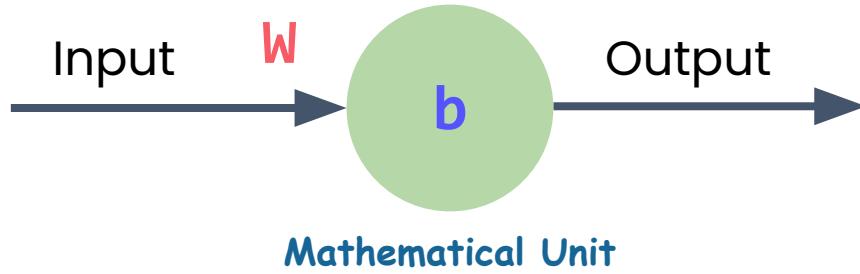
Prediction:

7 miles



$$time = w * distance + b$$

Single Neuron



$$output = W * input + b$$

Single Neuron

$$time = 1.0 * distance + 10.0$$

Distance (miles)	Delivery Time (minutes)
2.0	11.9
3.0	15.3
4.0	18.8
5.0	22.2
6.0	25.6



Prediction:

$$1.0 * 5 + 10.0 = 15$$

Single Neuron

$$time = 3.4 * distance + 5.0$$

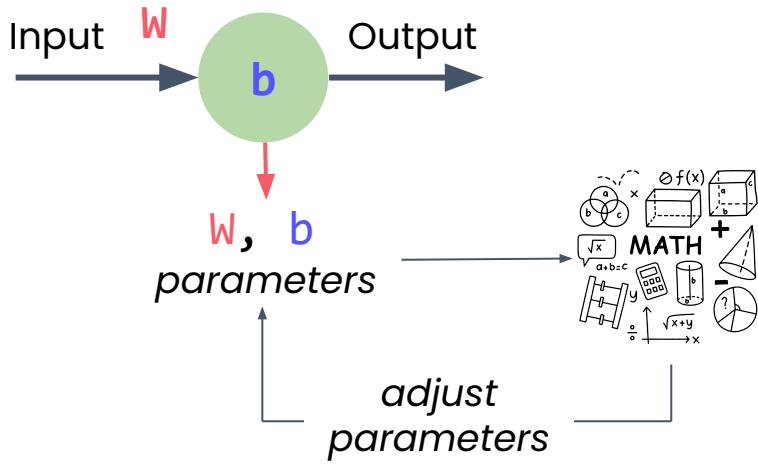
Distance (miles)	Delivery Time (minutes)
2.0	11.9
3.0	15.3
4.0	18.8
5.0	22.2
6.0	25.6



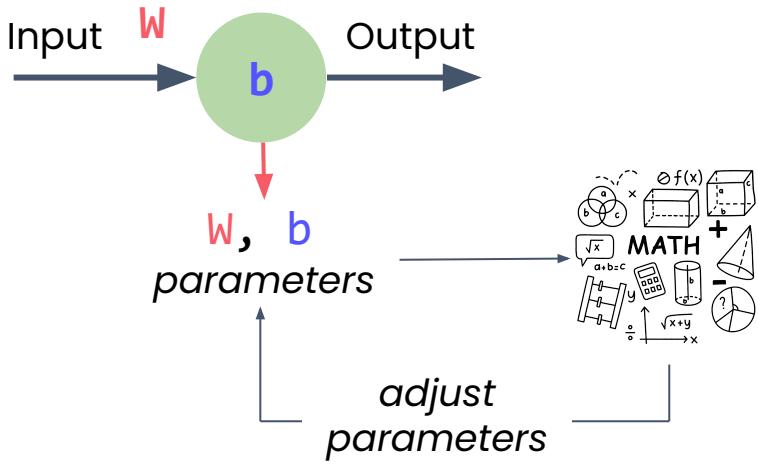
Prediction:

$$3.4 * 5 + 5.0 = 22$$

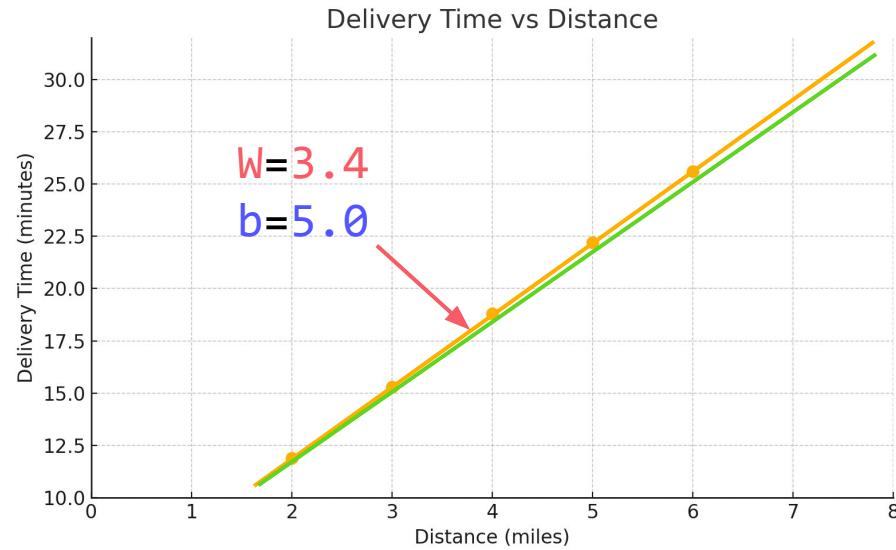
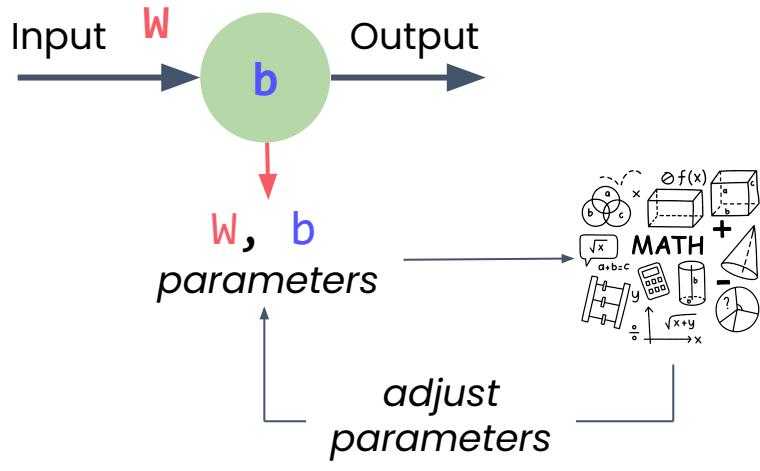
Single Neuron



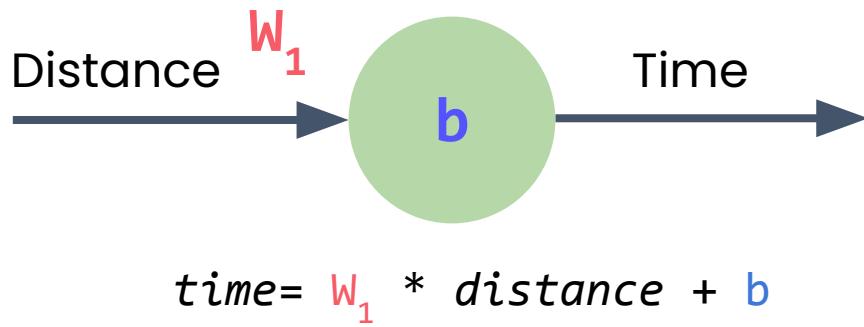
Single Neuron



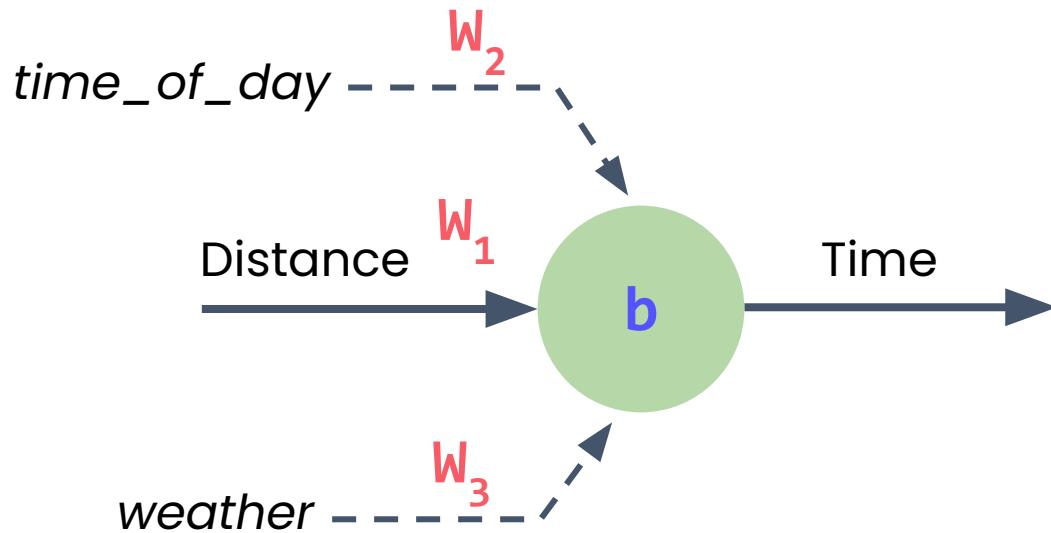
Single Neuron



Single Neuron

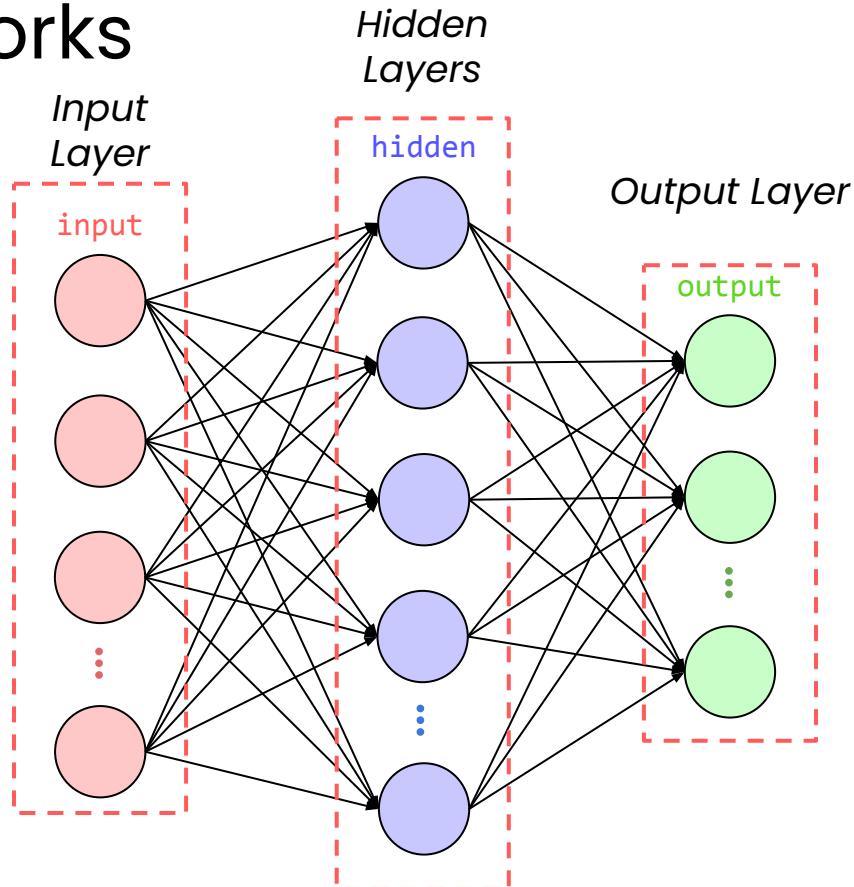


Single Neuron



$$time = w_1 * distance + w_2 * time_of_day + w_3 * weather + b$$

Neural Networks



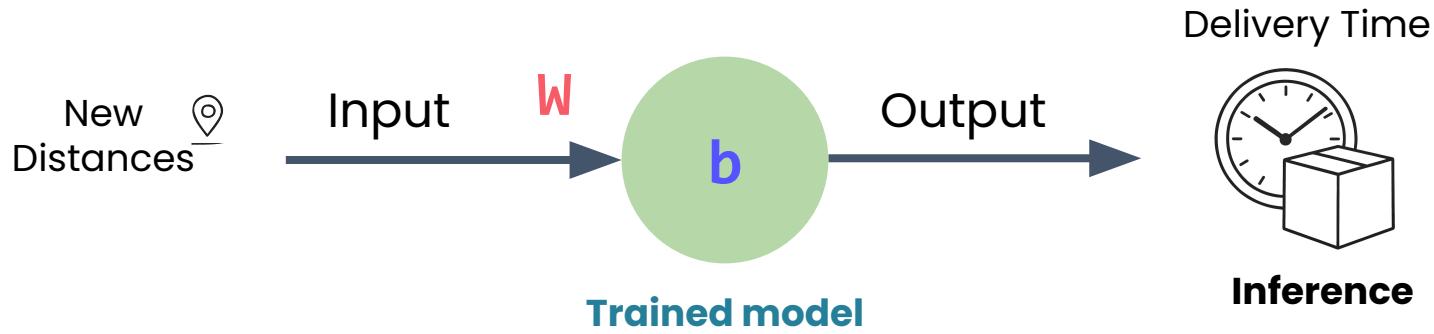


DeepLearning.AI

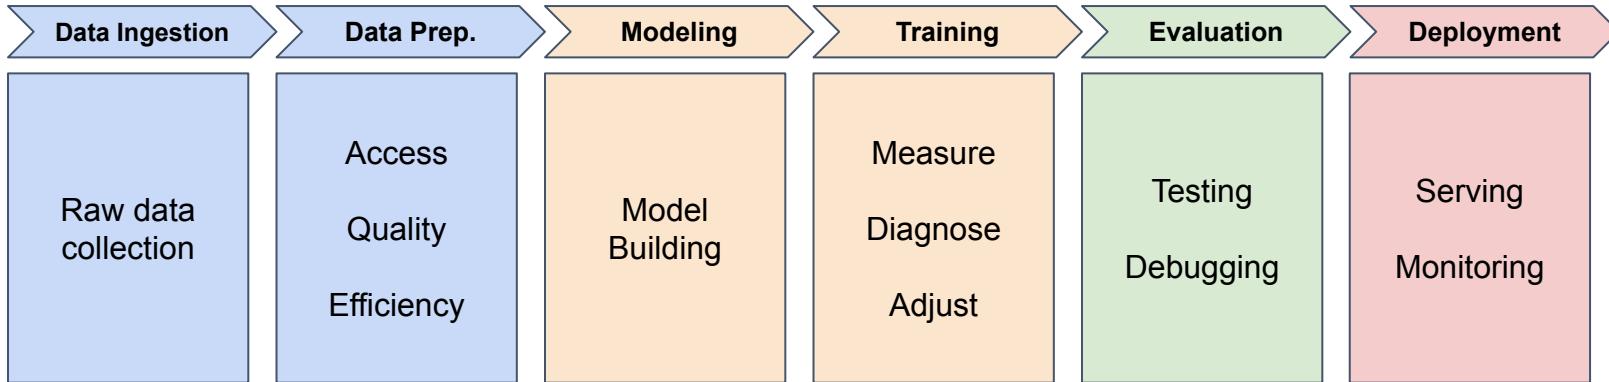
The ML Pipeline

Getting Started with PyTorch

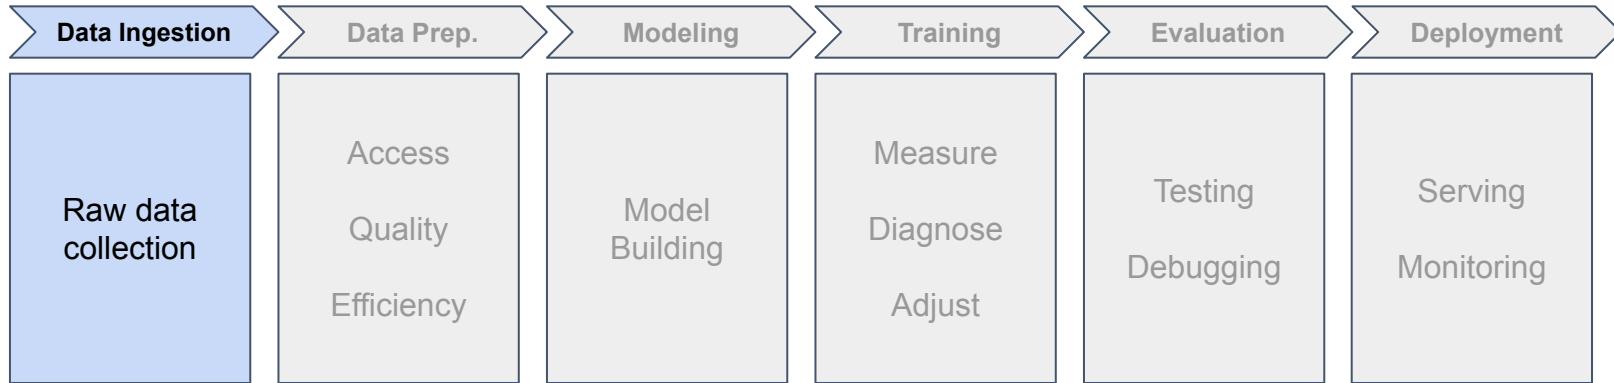
Inference



The ML pipeline



Data Ingestion



Data Ingestion



Distance (miles)	Delivery Time (minutes)
2.0	11.9
3.0	15.3
4.0	18.8
5.0	22.2
6.0	25.6

Data Ingestion



Distance (miles)	Delivery Time (minutes)
7	"22 minutes"

Data Ingestion



Distance (miles)	Delivery Time (minutes)
7	"22 minutes"
7	22.0

Data Ingestion



Distance (miles)	Delivery Time (minutes)
7	"22 minutes"
7	22.0
N/A	18.8

Data Ingestion



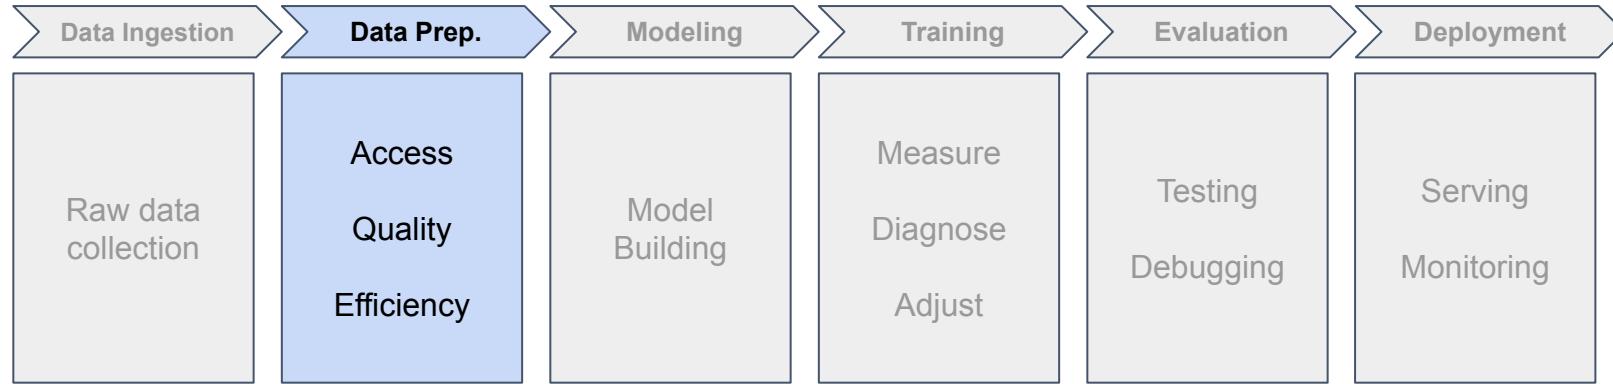
Distance (miles)	Delivery Time (minutes)
7	"22 minutes"
7	22.0
N/A	18.8
N/A	-22.2

Data Ingestion



Distance (miles)	Delivery Time (minutes)
7	"22 minutes"
7	22.0
N/A	18.8
N/A	-22.2
10 miles	3 minutes

Data Preparation

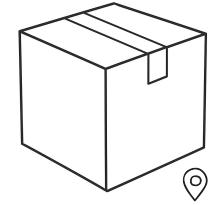


Data Preparation

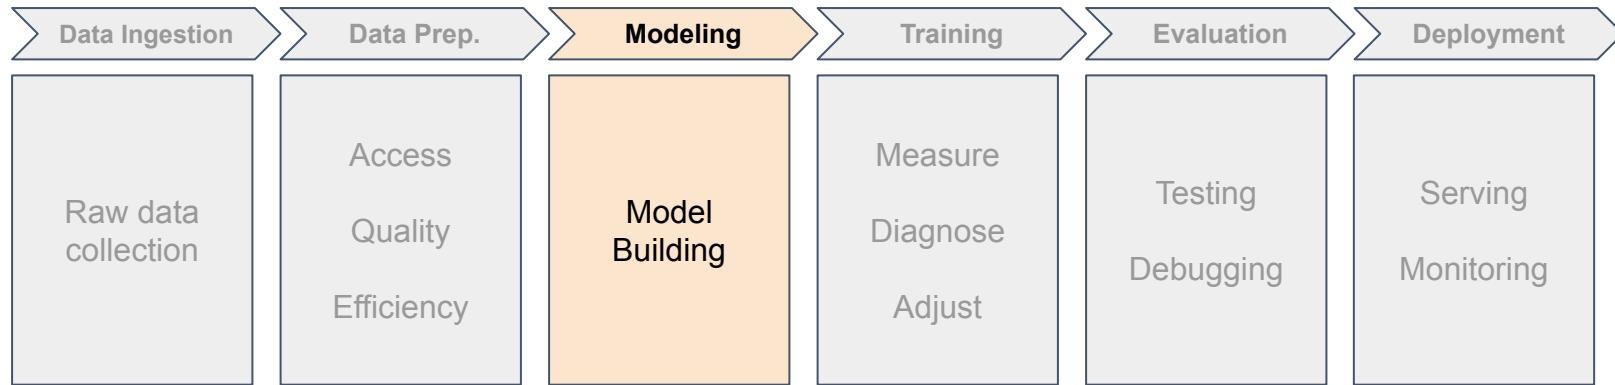


Sources of data

-  delivery records
 - Cleaning
 - Fixing errors
-  customer images →
 - Transforming →
 - Handling missing values
-  email text
 - Organizing
 - Engineering new features
-  other



Model Building



Model Building

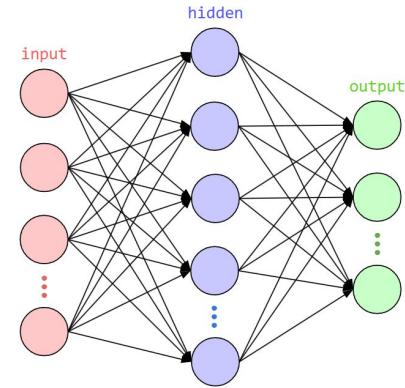


Output

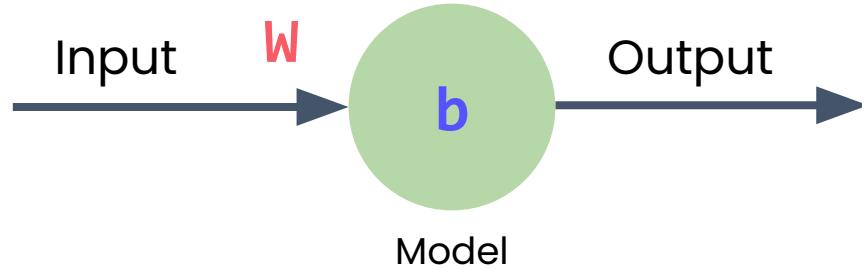
- Delivery times
- Classifying images
- Analyzing text
- Others

Choosing the architecture

- How many neurons should you use?
- How are they connected?
- What kinds of layers does the model need?
- Cost, time, infrastructure, others?

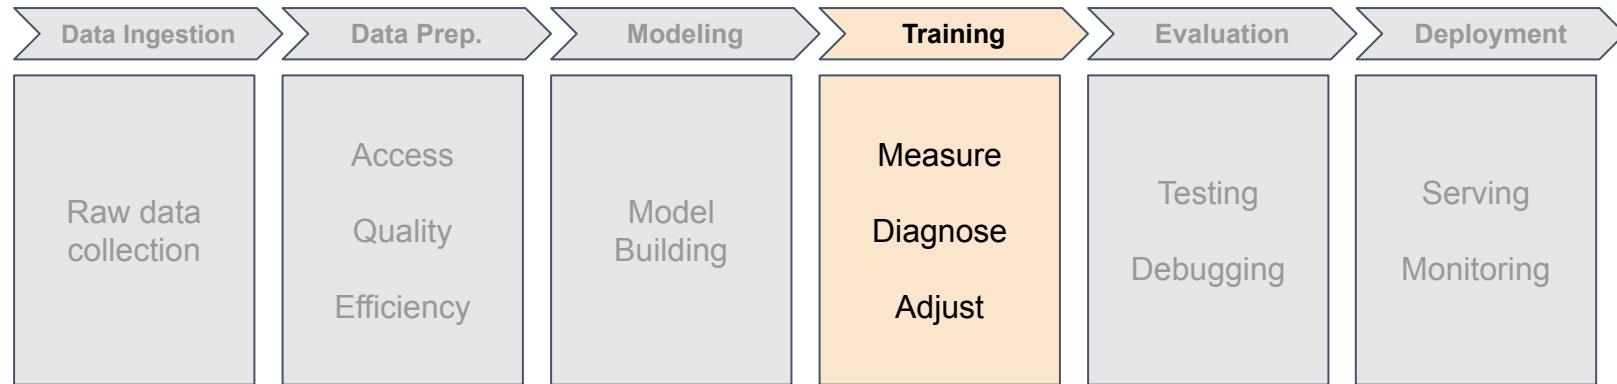


Model Building

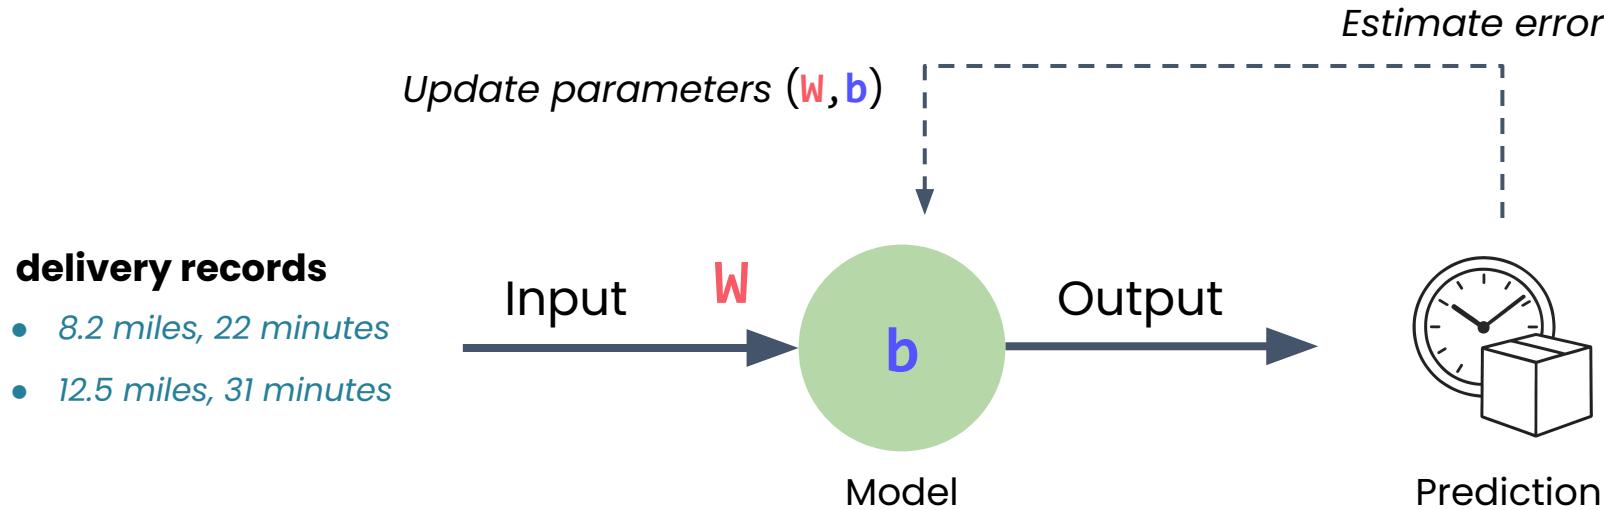


```
model = nn.Sequential(nn.Linear(1, 1))
```

Training



Training



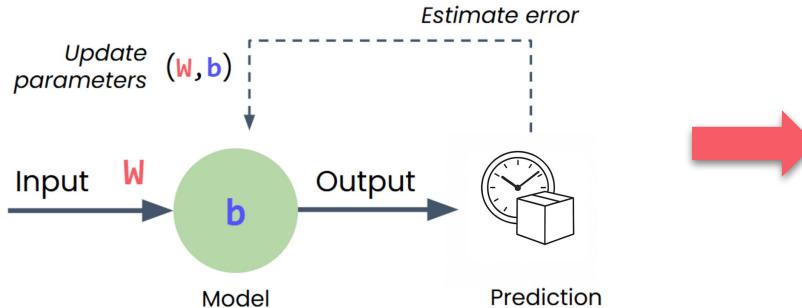
Training



Configure

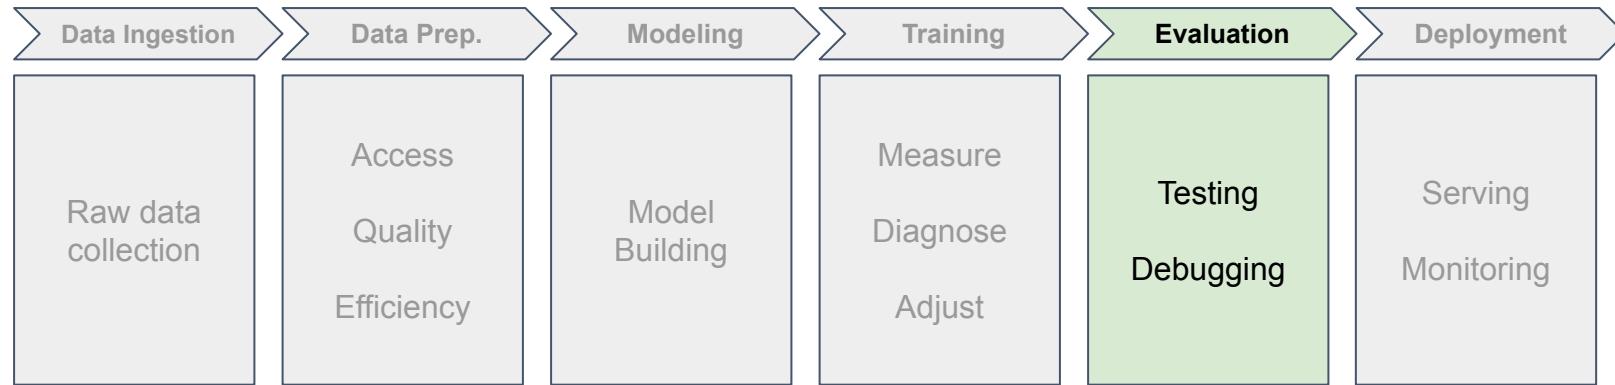
- How to measure prediction errors
- How to guide the model to improve
- How fast it learns

Training Loop

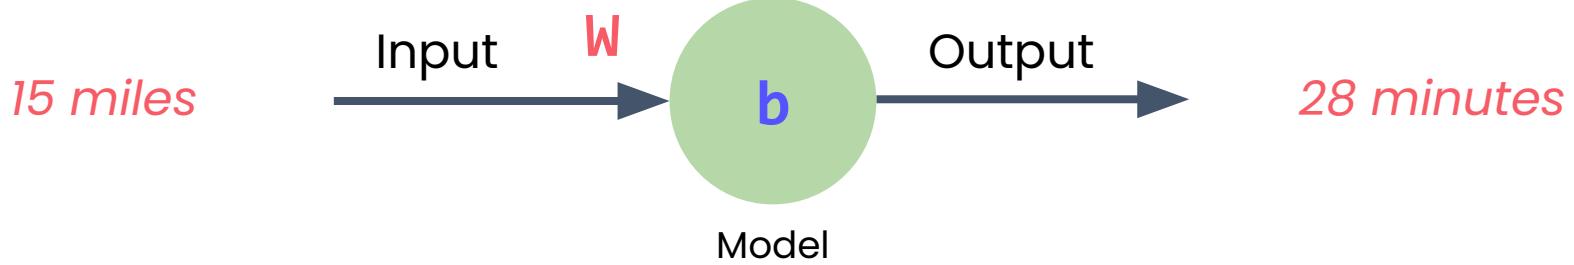


Good
predictions?

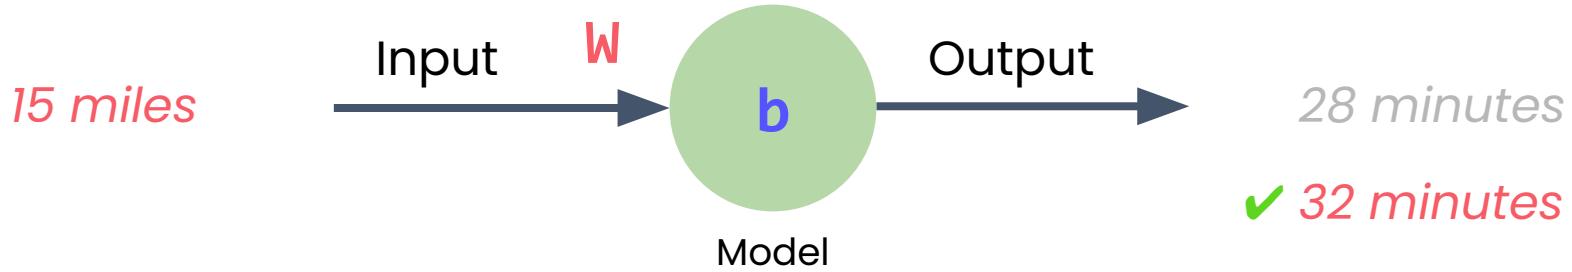
Evaluation



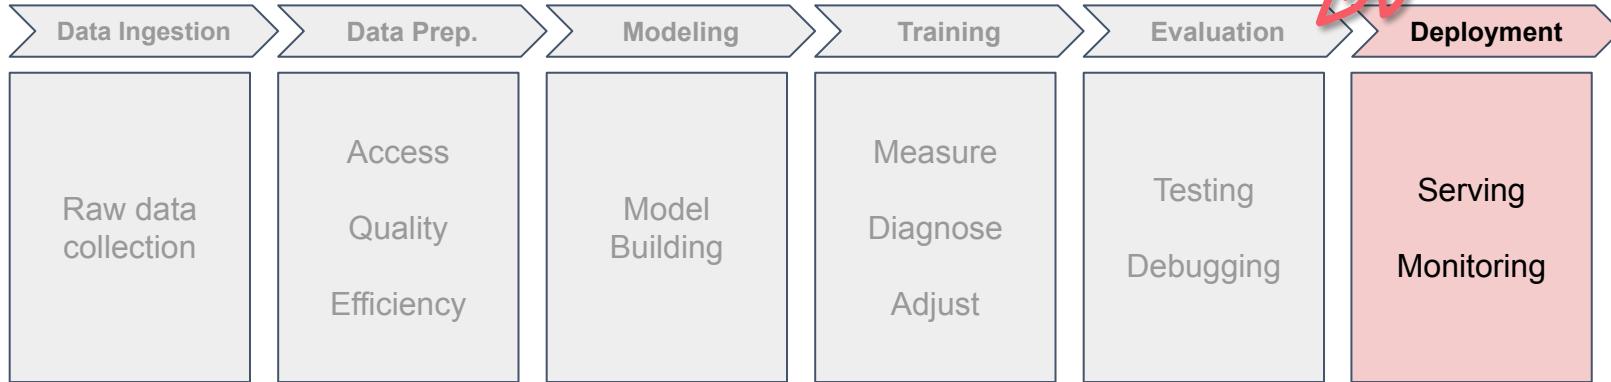
Evaluation & Debugging



Evaluation & Debugging



Deployment





DeepLearning.AI

Building a Simple Neural Network

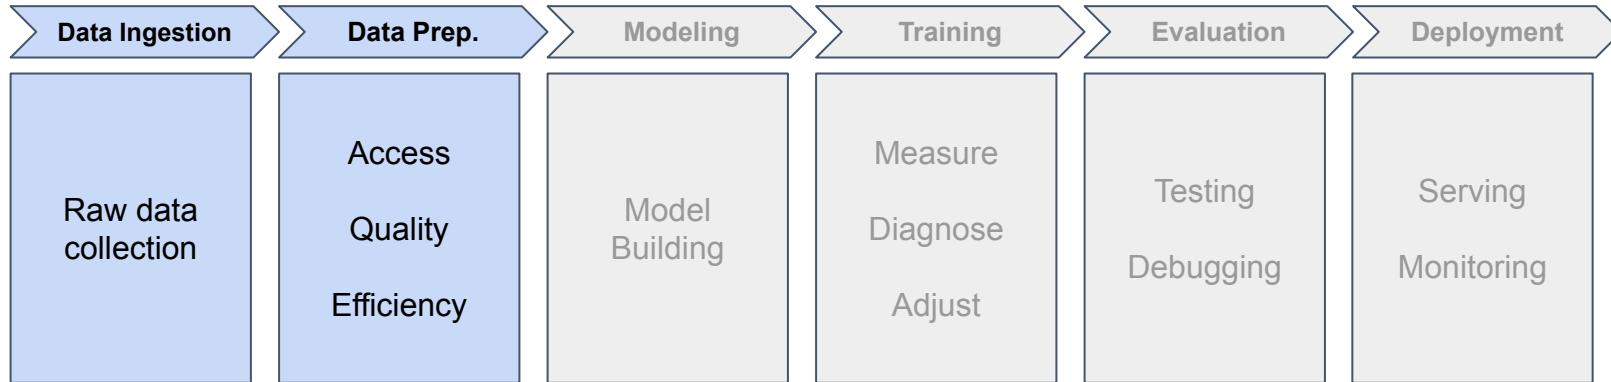
Getting Started with PyTorch

Import Libraries

```
import torch
import torch.nn as nn
import torch.optim as optim
```

- PyTorch's core functionality -> `torch`
- Components for building neural networks -> `nn`
- And tools for training those networks -> `optim`

The ML Pipeline



Defining Data



```
# Distance in miles
distances = torch.tensor([[1.0],[2.0],[3.0],[4.0]], dtype=torch.float32)
# Delivery times in minutes
times = torch.tensor([[6.96],[12.11],[16.77],[22.21]], dtype=torch.float32)
```

Defining Data



```
# Distance in miles  
distances = torch.tensor([[1.0],[2.0],[3.0],[4.0]], dtype=torch.float32)  
# Delivery times in minutes  
times = torch.tensor([[6.96],[12.11],[16.77],[22.21]], dtype=torch.float32)
```

Defining Data



```
# Distance in miles  
distances = torch.tensor([[1.0],[2.0],[3.0],[4.0]], dtype=torch.float32)  
# Delivery times in minutes  
times = torch.tensor([[6.96],[12.11],[16.77],[22.21]], dtype=torch.float32)
```

Defining Data



```
# Distance in miles  
distances = torch.tensor([[1.0],[2.0],[3.0],[4.0]], dtype=torch.float32)  
# Delivery times in minutes  
times = torch.tensor([[6.96],[12.11],[16.77],[22.21]], dtype=torch.float32)
```

Defining Data - Multiple Inputs



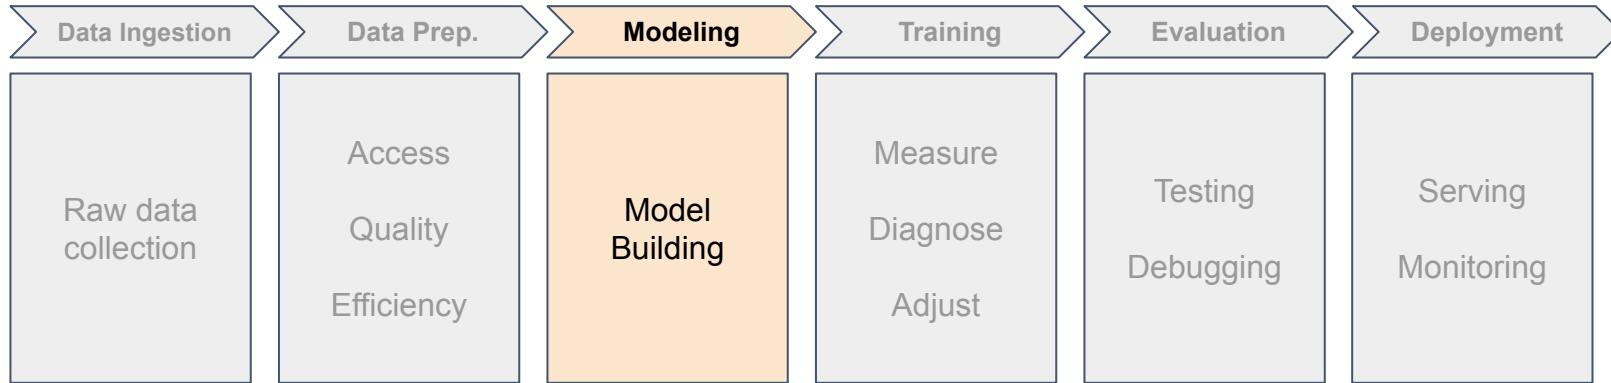
```
# Distance in miles  
inputs = torch.tensor([[2.0,3.0,4.0],[6.0,7.0,8.0],[11.0,12.0,13.0],...],dtype=torch.float32)  
# Delivery times in minutes  
times = torch.tensor([[14.8],[24.1],[36.5],[51.2],[60.8],[75.3]],dtype=torch.float32)
```

Defining Data



```
# Distance in miles  
distances = torch.tensor([[1.0],[2.0],[3.0],[4.0]]), dtype=torch.float32  
# Delivery times in minutes  
times = torch.tensor([[6.96],[12.11],[16.77],[22.21]]), dtype=torch.float32
```

Model Building



Model Building



```
# Define the model
model = nn.Sequential(nn.Linear(1, 1))
```

Model Building

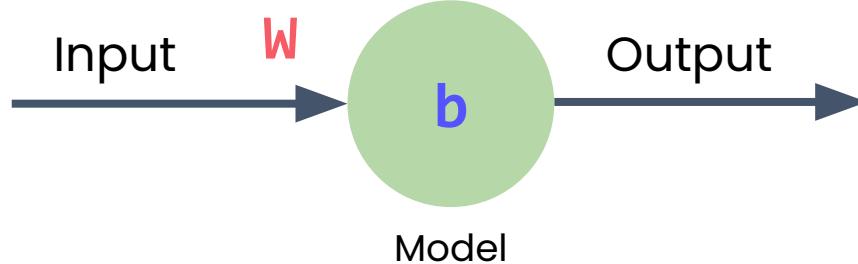


```
# Define the model
model = nn.Sequential(nn.Linear(1, 1))
```

Model Building



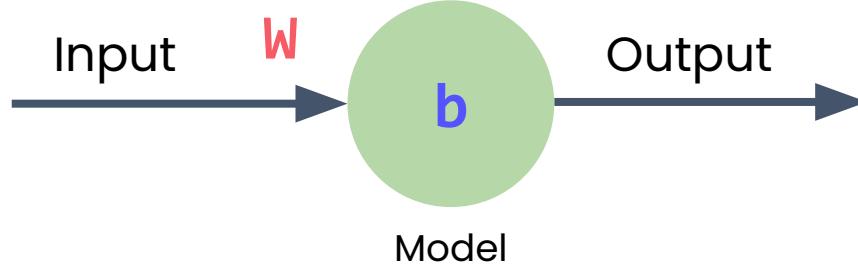
```
# Define the model  
model = nn.Sequential(nn.Linear(1, 1))
```



Model Building



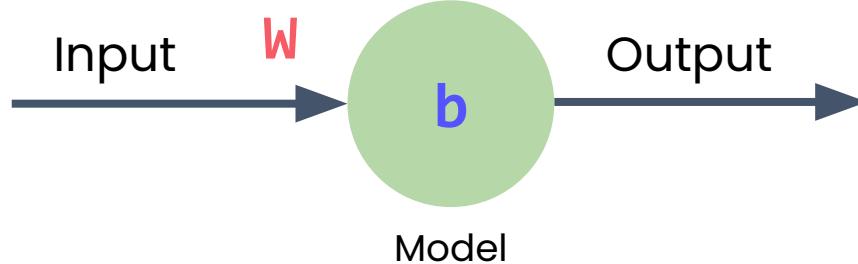
```
# Define the model  
model = nn.Sequential(nn.Linear(1, 1))
```



Model Building



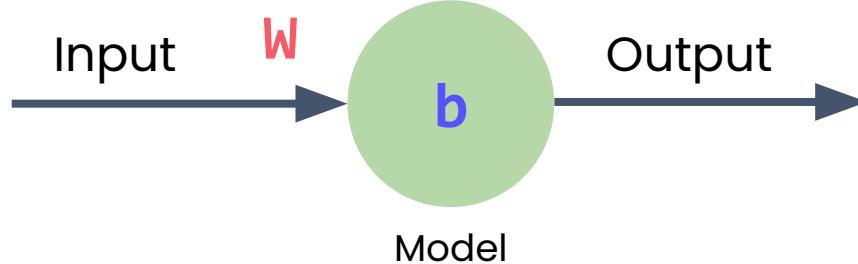
```
# Define the model  
model = nn.Sequential(nn.Linear(1, 1))
```



Model Building



```
# Define the model  
model = nn.Sequential(nn.Linear(1, 1))
```

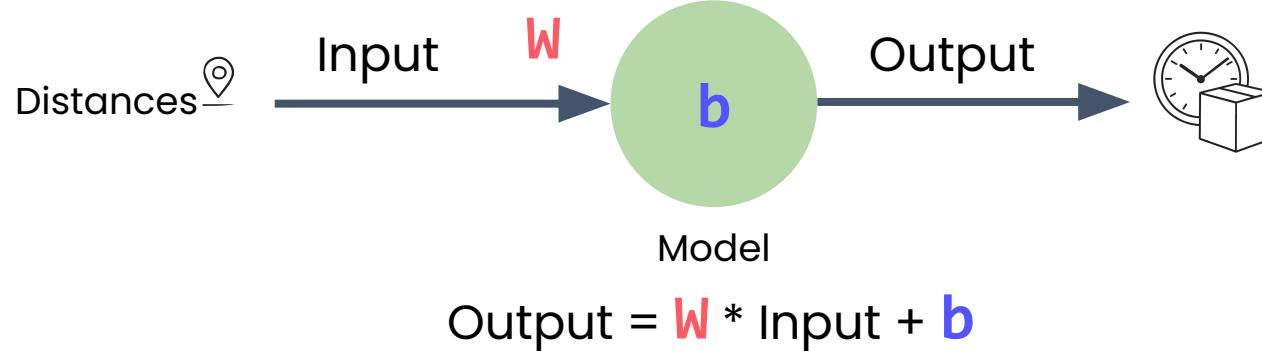


$$\text{Output} = W * \text{Input} + b$$

Model Building



```
# Define the model  
model = nn.Sequential(nn.Linear(1, 1))
```



Loss Function & Optimizer



```
# Define the loss function and optimizer
loss_function = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)
```

Loss Function & Optimizer



```
# Define the loss function and optimizer
loss_function = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)
```

Mean Squared Error Loss

Loss Function & Optimizer



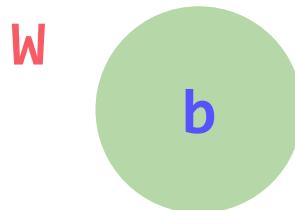
```
# Define the loss function and optimizer
loss_function = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)
```

Stochastic Gradient Descent

Loss Function & Optimizer



```
# Define the loss function and optimizer
loss_function = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)
```



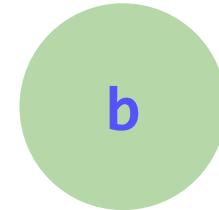
$$time = W * distance + b$$

Loss Function & Optimizer



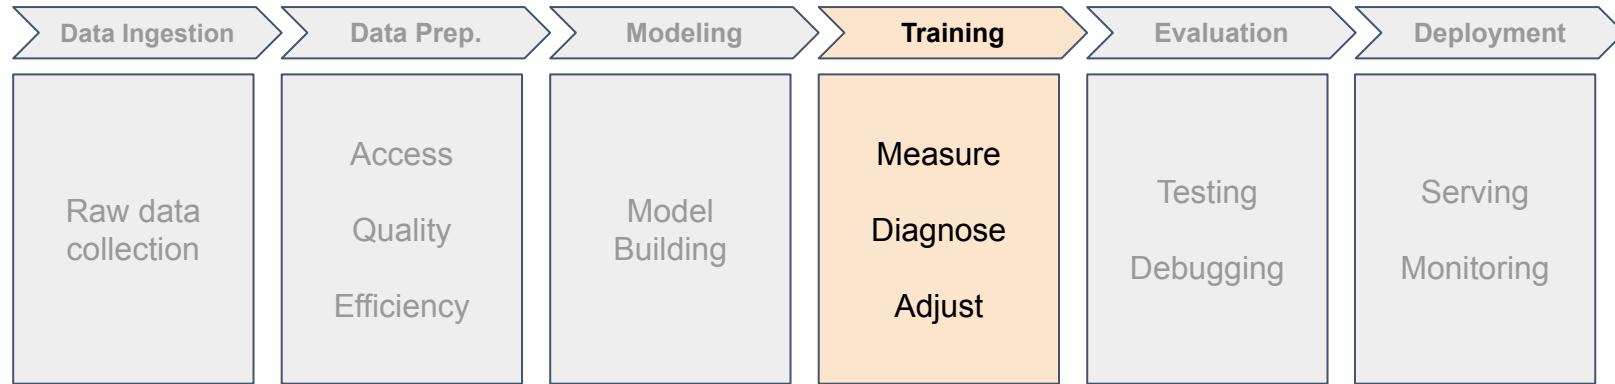
```
# Define the loss function and optimizer
loss_function = nn.MSELoss()
optimizer = optim.SGD(model.parameters(), lr=0.01)
```

Learning Rate



$$time = W * distance + b$$

Training



Training



```
# Training loop
for epoch in range(500):
    # 0. Reset the optimizer
    optimizer.zero_grad()
    # 1. Make predictions
    outputs = model(inputs)
    # 2. Calculate the loss - how bad was this guess?
    loss = loss_function(outputs, times)
    # 3. Calculate adjustments
    loss.backward()
    # 4. Update the model
    optimizer.step()
```

Training



```
# Training loop
for epoch in range(500):
    # 0. Reset the optimizer
    optimizer.zero_grad()
    # 1. Make predictions
    outputs = model(inputs)
    # 2. Calculate the loss - how bad was this guess?
    loss = loss_function(outputs, times)
    # 3. Calculate adjustments
    loss.backward()
    # 4. Update the model
    optimizer.step()
```

Training



```
# Training loop
for epoch in range(500):
    # 0. Reset the optimizer
    optimizer.zero_grad()
    # 1. Make predictions
    outputs = model(inputs)
    # 2. Calculate the loss - how bad was this guess?
    loss = loss_function(outputs, times)
    # 3. Calculate adjustments
    loss.backward()
    # 4. Update the model
    optimizer.step()
```

Training



```
# Training loop
for epoch in range(500):
    # 0. Reset the optimizer
    optimizer.zero_grad()
    # 1. Make predictions
    outputs = model(inputs)
    # 2. Calculate the loss - how bad was this guess?
    loss = loss_function(outputs, times)
    # 3. Calculate adjustments
    loss.backward()
    # 4. Update the model
    optimizer.step()
```

Training



```
# Training loop
for epoch in range(500):
    # 0. Reset the optimizer
    optimizer.zero_grad()
    # 1. Make predictions
    outputs = model(inputs)
    # 2. Calculate the loss - how bad was this guess?
    loss = loss_function(outputs, times)
    # 3. Calculate adjustments
    loss.backward()
    # 4. Update the model
    optimizer.step()
```

Training



```
# Training loop
for epoch in range(500):
    # 0. Reset the optimizer
    optimizer.zero_grad()
    # 1. Make predictions
    outputs = model(inputs)
    # 2. Calculate the loss - how bad was this guess?
    loss = loss_function(outputs, times)
    # 3. Calculate adjustments
    loss.backward()
    # 4. Update the model
    optimizer.step()
```

Training



```
# Training loop
for epoch in range(500):
    # 0. Reset the optimizer
    optimizer.zero_grad()
    # 1. Make predictions
    outputs = model(inputs)
    # 2. Calculate the loss - how bad was this guess?
    loss = loss_function(outputs, times)
    # 3. Calculate adjustments
    loss.backward()
    # 4. Update the model
    optimizer.step()
```

Training



```
# Training loop
for epoch in range(500):
    # 0. Reset the optimizer
    optimizer.zero_grad()
    # 1. Make predictions
    outputs = model(inputs)
    # 2. Calculate the loss - how bad was this guess?
    loss = loss_function(outputs, times)
    # 3. Calculate adjustments
    loss.backward()
    # 4. Update the model
    optimizer.step()
```

Training



```
# Training loop
for epoch in range(500):
    # 0. Reset the optimizer
    optimizer.zero_grad()
    # 1. Make predictions
    outputs = model(inputs)
    # 2. Calculate the loss - how bad was this guess?
    loss = loss_function(outputs, times)
    # 3. Calculate adjustments
    loss.backward()    Back Propagation
    # 4. Update the model
    optimizer.step()
```

Training



```
# Training loop
for epoch in range(500):
    # 0. Reset the optimizer
    optimizer.zero_grad()
    # 1. Make predictions
    outputs = model(inputs)
    # 2. Calculate the loss - how bad was this guess?
    loss = loss_function(outputs, times)
    # 3. Calculate adjustments
    loss.backward()
    # 4. Update the model
    optimizer.step()
```

Inference

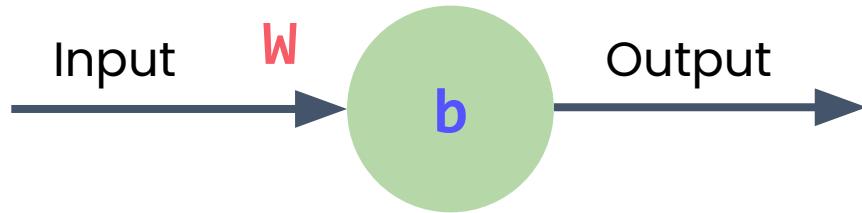
```
with torch.no_grad():
    test_distance = torch.tensor([[25.0]], dtype=torch.float32)
    predicted_time = model(test_distance)
    print(f"Predicted time for 25 miles: {predicted_time.item():.1f} minutes")
```

Inference

```
with torch.no_grad():
    test_distance = torch.tensor([[25.0]], dtype=torch.float32)
    predicted_time = model(test_distance)
    print(f"Predicted time for 25 miles: {predicted_time.item():.1f} minutes")
```

Inference

```
with torch.no_grad():
    test_distance = torch.tensor([[25.0]], dtype=torch.float32)
    predicted_time = model(test_distance)
    print(f"Predicted time for 25 miles: {predicted_time.item():.1f} minutes")
```





DeepLearning.AI

Activation Functions

Getting Started with PyTorch

The Delivery Problem



The Delivery Problem



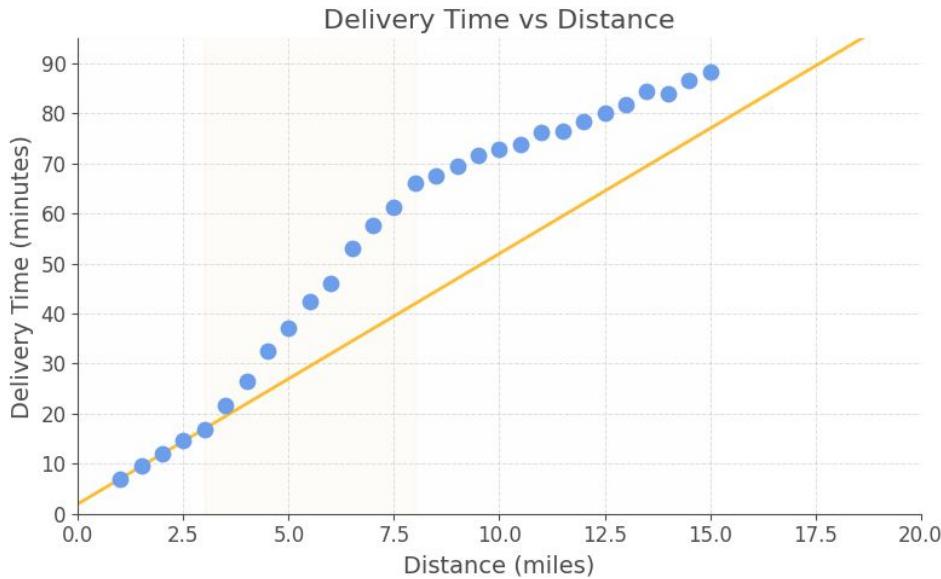
The Delivery Problem



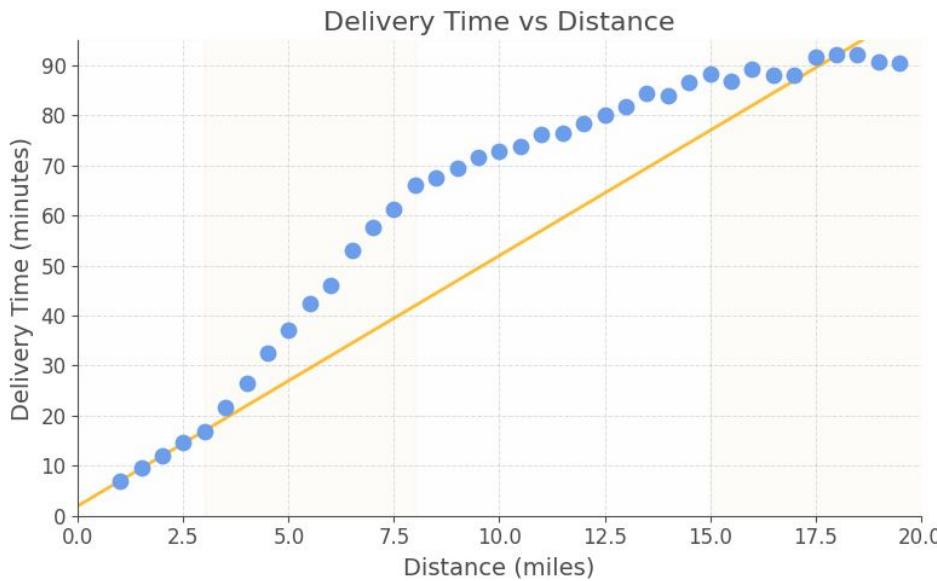
The Delivery Problem



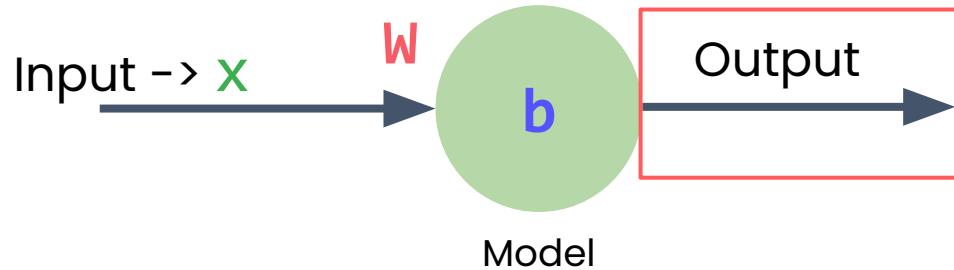
The Delivery Problem



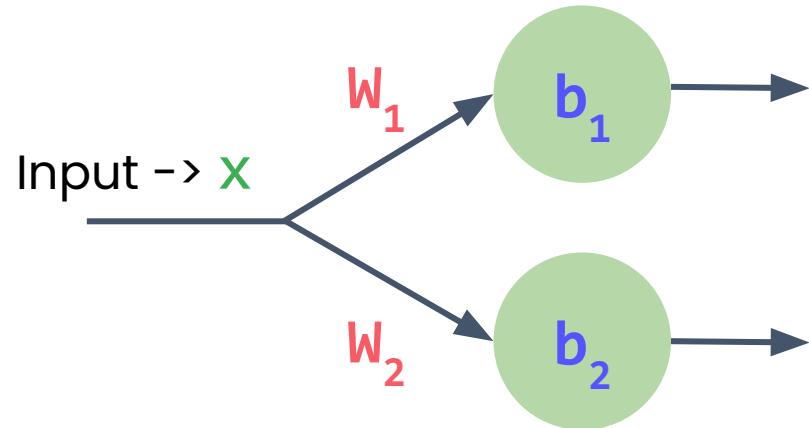
The Delivery Problem



Single Neurons

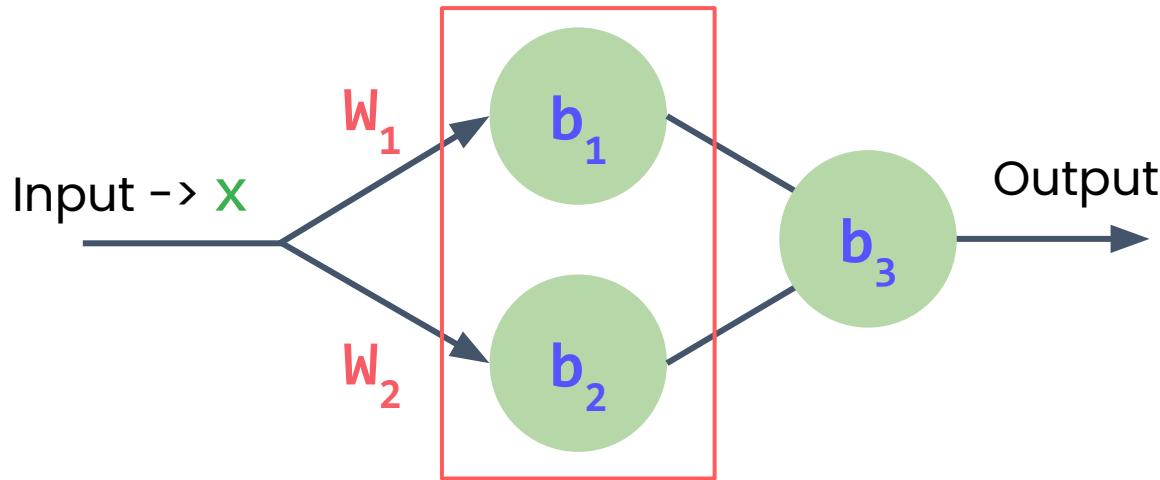


Two Neurons

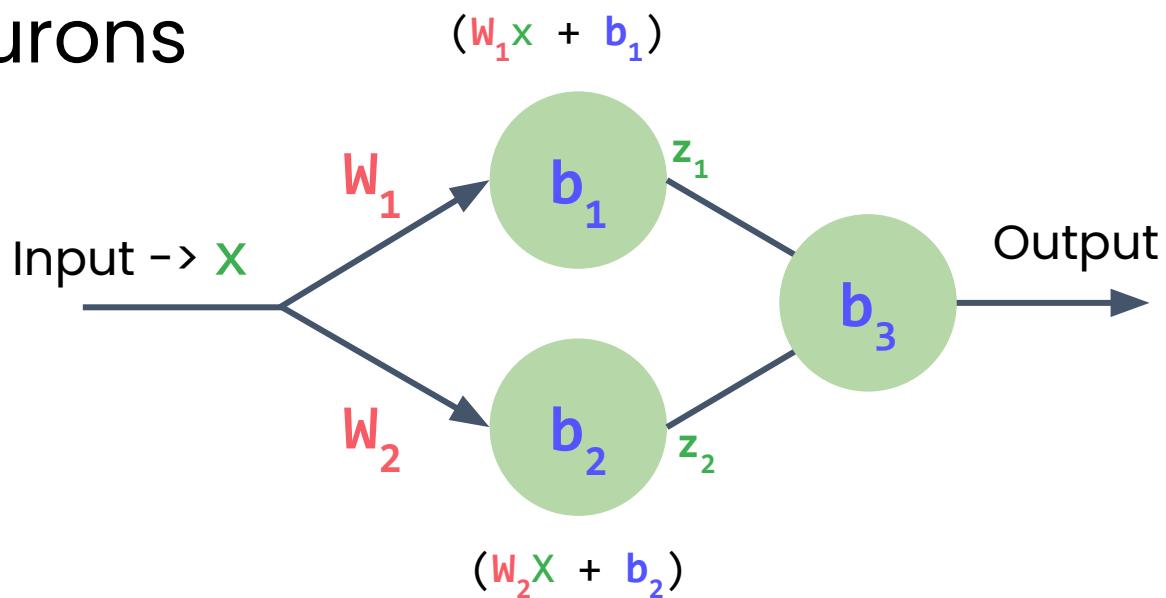


Two Neurons

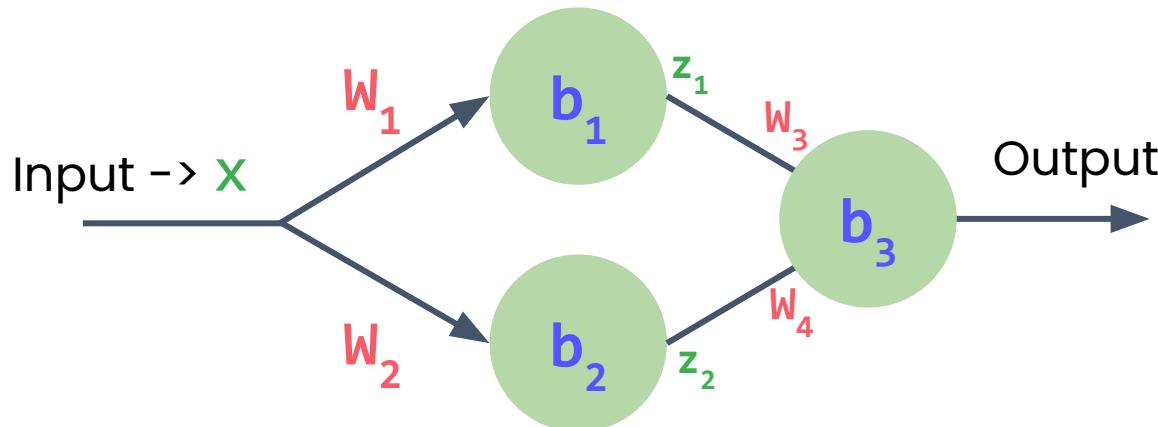
Hidden Layer



Two Neurons



Two Neurons



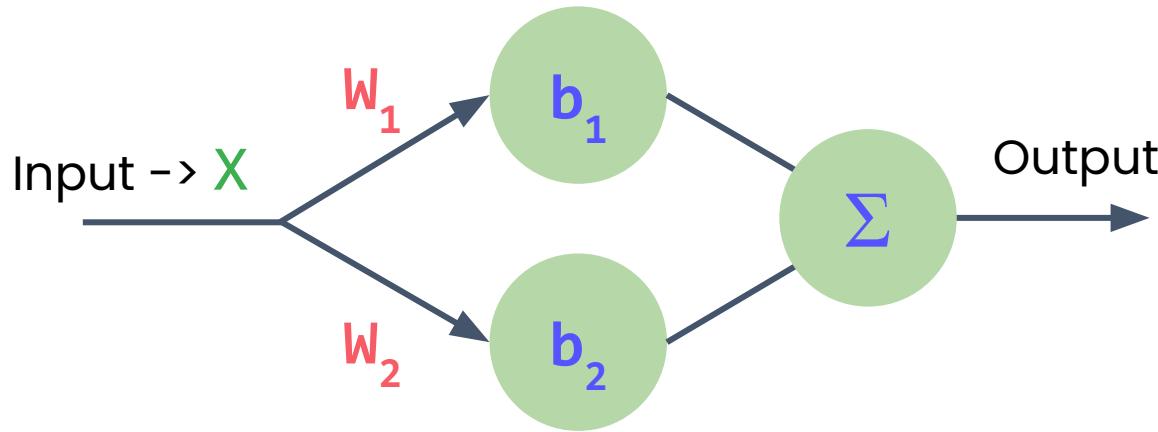
$$\text{Output} = w_3(w_1x + b_1) + w_4(w_2x + b_2) + b_3$$

$$(w_3w_1 + w_4w_2)x + (w_3b_1 + w_4b_2 + b_3)$$

$$w^*x$$

Still a linear equation!

Two Neurons

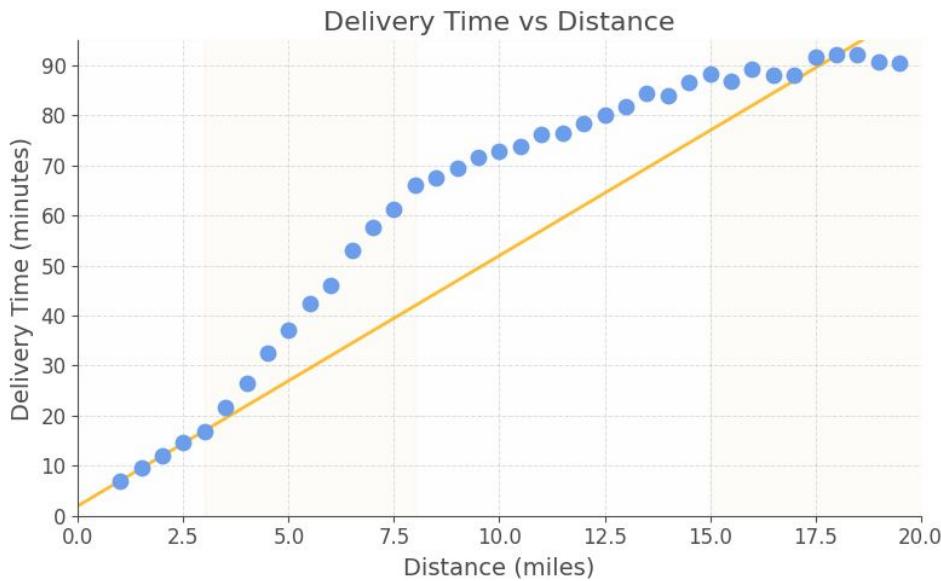


$$\text{Output} = (w_1 X + b_1) + (w_2 X + b_2)$$

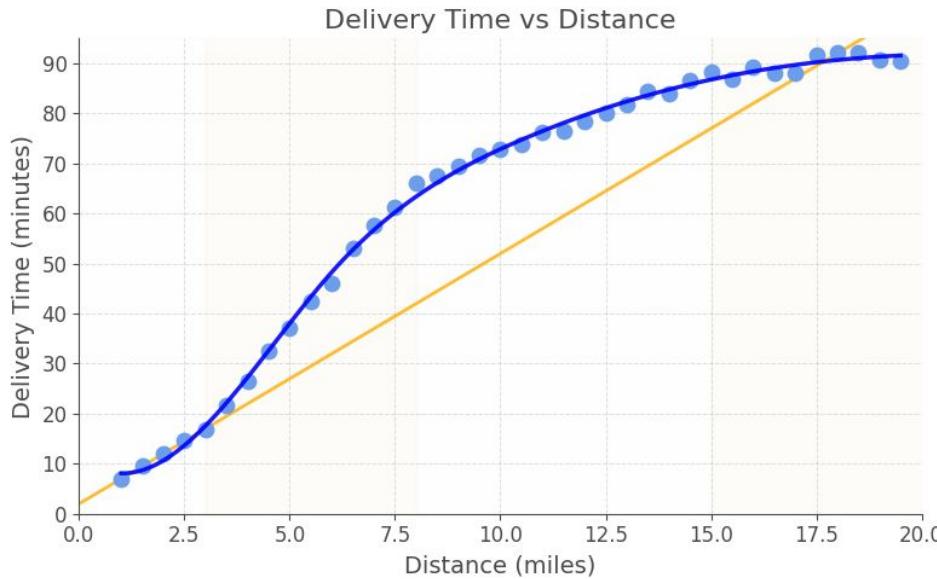
$$(w_1 + w_2) * X + (b_1 + b_2)$$

$w * X + b$ Still a linear equation!

The Delivery Problem

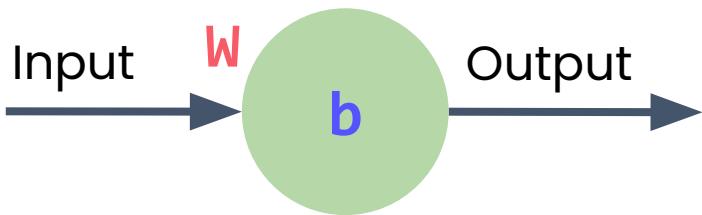


The Delivery Problem

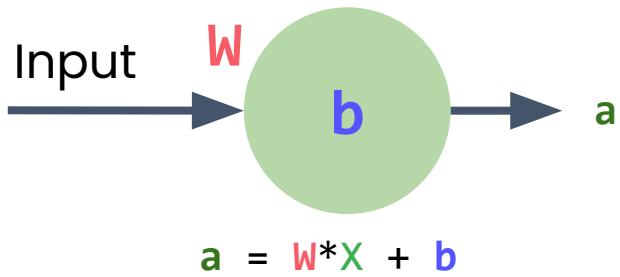


$$w^*x + b$$

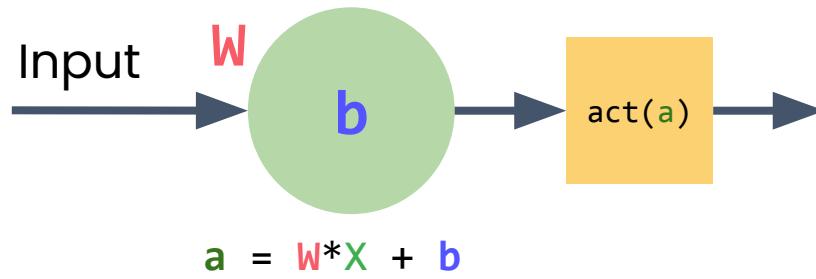
Activation Functions



Activation Functions



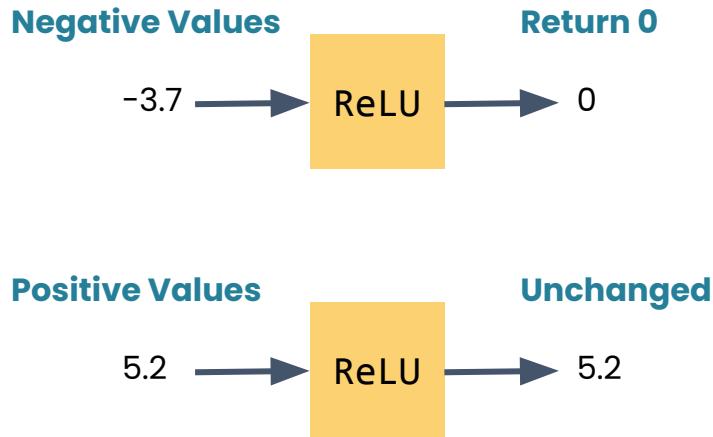
Activation Functions



$$\text{Output} = \text{act}(a)$$

ReLU

`nn.ReLU()`



```
if input > 0:  
    return input  
else:  
    return 0
```

ReLU

```
# Define the model
model = nn.Sequential(
    nn.Linear(1, 1),
    nn.ReLU()
)
```

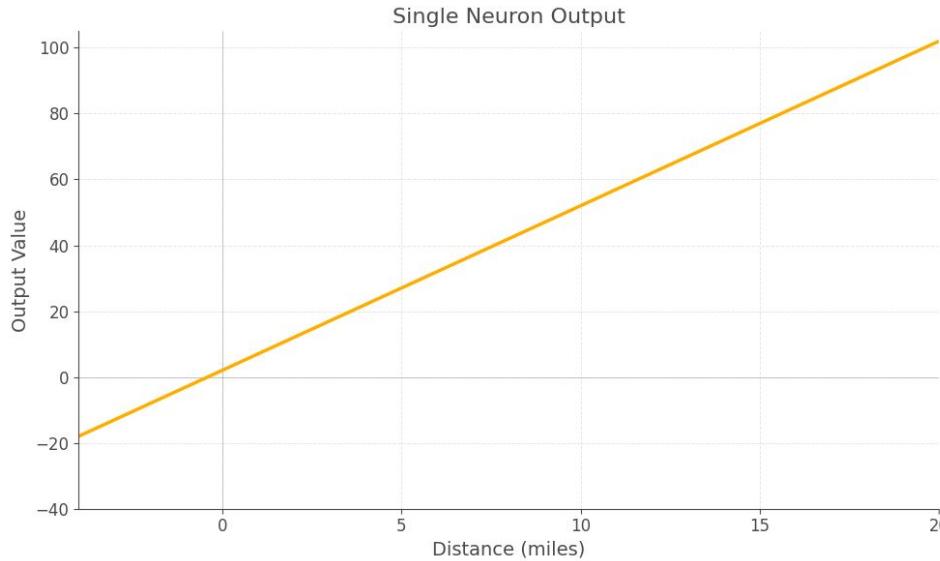
ReLU

```
# Define the model
model = nn.Sequential(
    nn.Linear(1, 1),
    nn.ReLU()
)
```

ReLU

```
# Define the model
model = nn.Sequential(
    nn.Linear(1, 1),
    nn.ReLU()
)
```

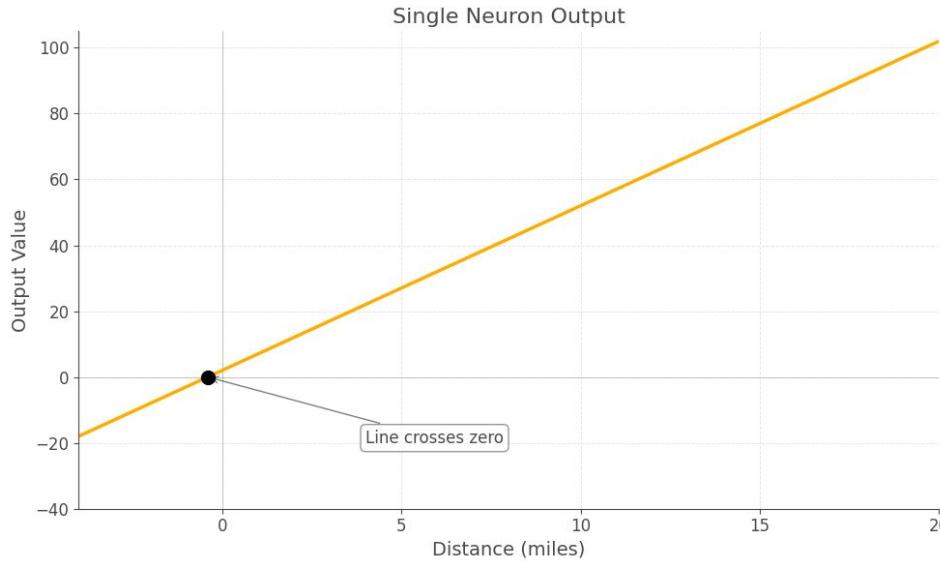
ReLU



$$time = w * distance + b$$

$$\text{ReLU}(time)$$

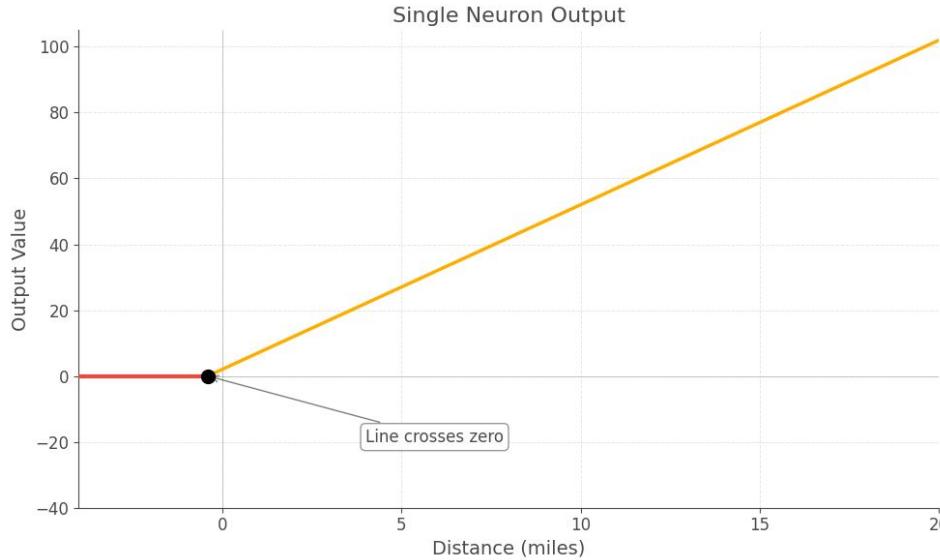
ReLU



$$time = w * distance + b$$

$$\text{ReLU}(time)$$

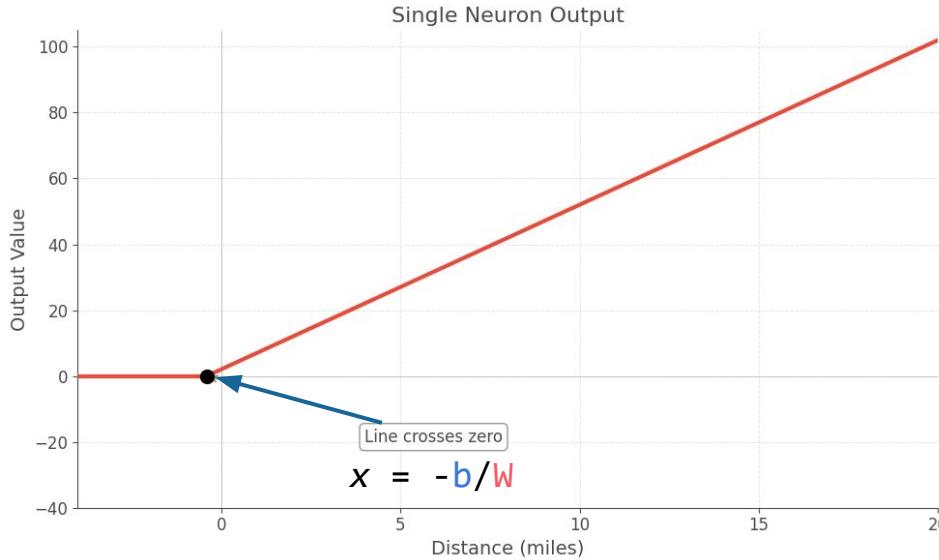
ReLU



$$time = w * distance + b$$

$$\text{ReLU}(time)$$

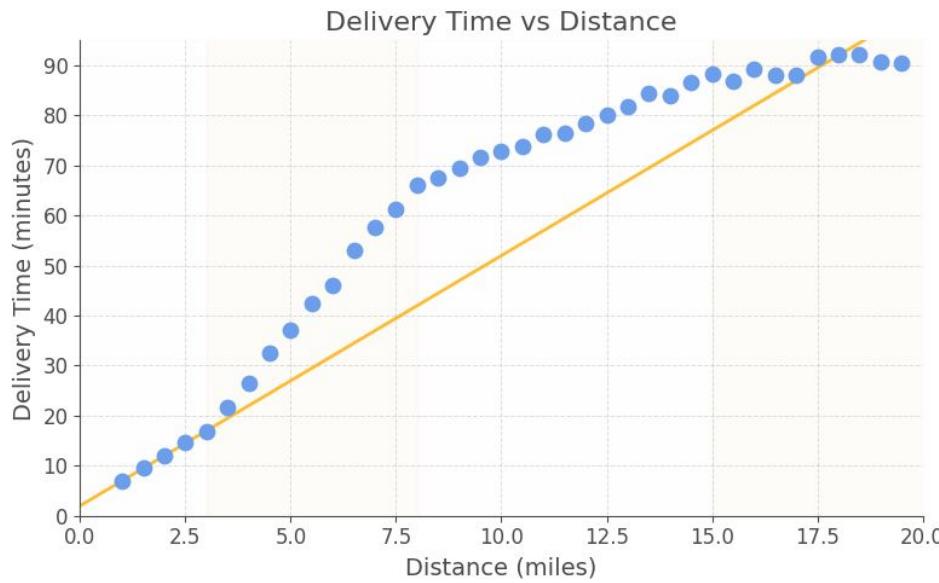
ReLU



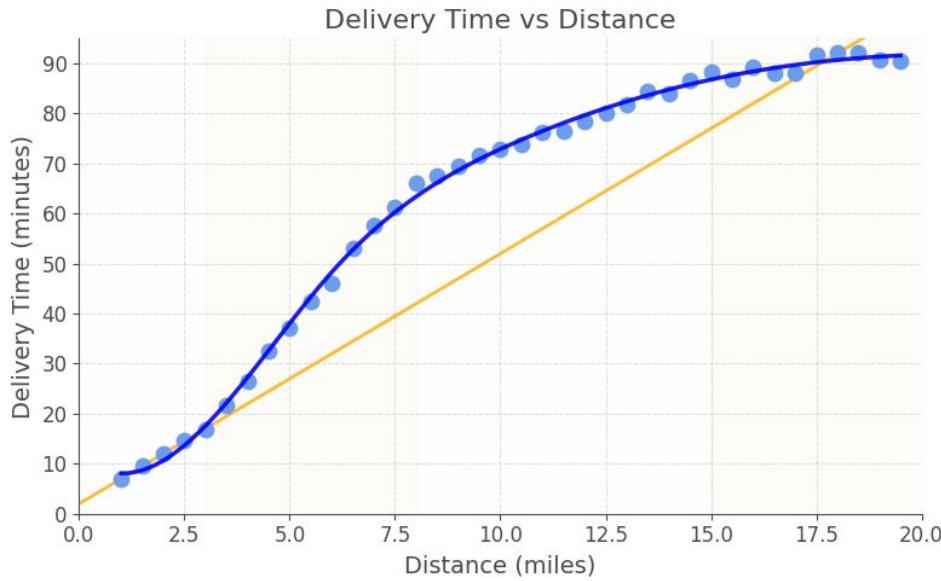
$$time = w * distance + b$$

$$\text{ReLU}(time)$$

The Delivery Problem



The Delivery Problem



ReLU



$$w_1 \ b_1$$

$$w_2 \ b_2$$

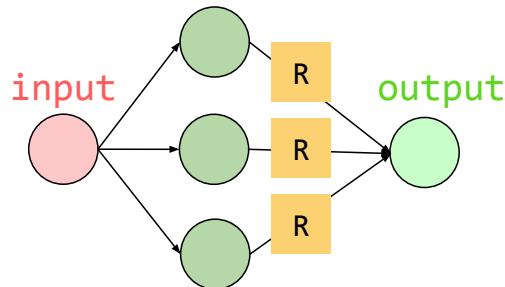
$$w_3 \ b_3$$

$$-b_1/w_1$$

$$-b_2/w_2$$

$$-b_3/w_3$$

ReLU

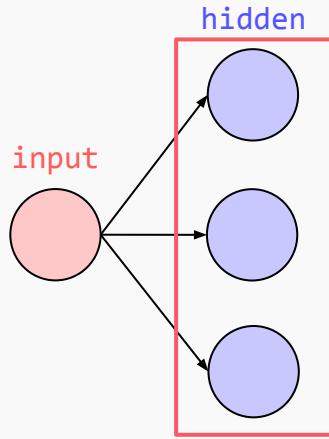


Non Linear Model

```
model = nn.Sequential(  
    nn.Linear(1, 3),  
    nn.ReLU(),  
    nn.Linear(3, 1)  
)
```

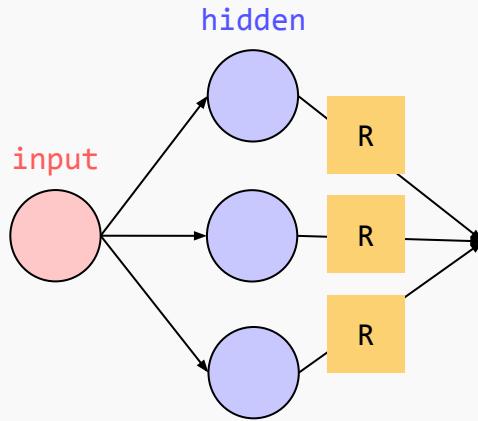
Non Linear Model

```
model = nn.Sequential(  
    nn.Linear(1, 3),  
    nn.ReLU(),  
    nn.Linear(3, 1)  
)
```



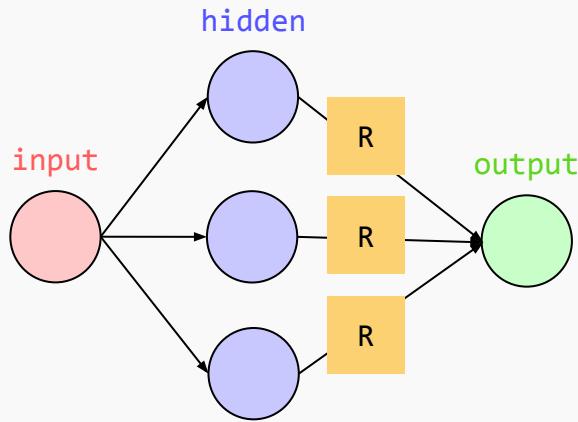
Non Linear Model

```
model = nn.Sequential(  
    nn.Linear(1, 3),  
    nn.ReLU(),  
    nn.Linear(3, 1)  
)
```



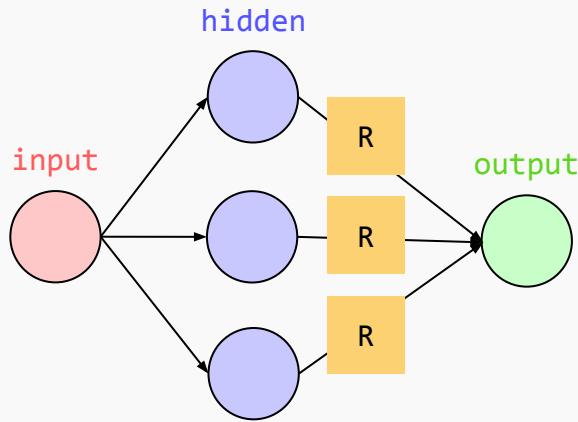
Non Linear Model

```
model = nn.Sequential(  
    nn.Linear(1, 3),  
    nn.ReLU(),  
    nn.Linear(3, 1)  
)
```



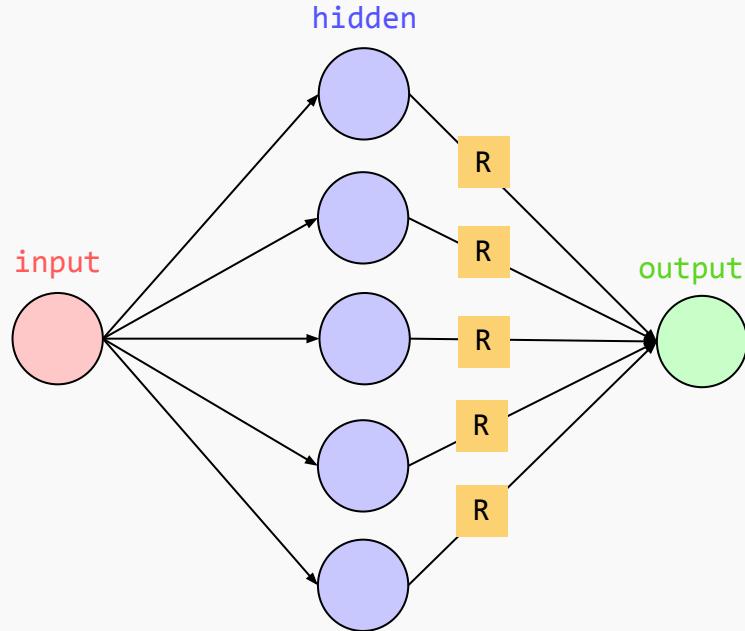
Non Linear Model

```
model = nn.Sequential(  
    nn.Linear(1, 3),  
    nn.ReLU(),  
    nn.Linear(3, 1)  
)
```

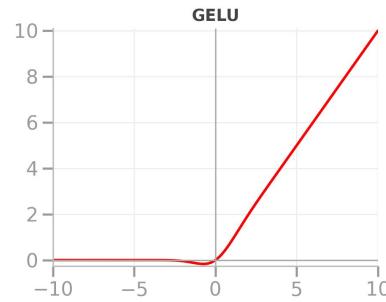
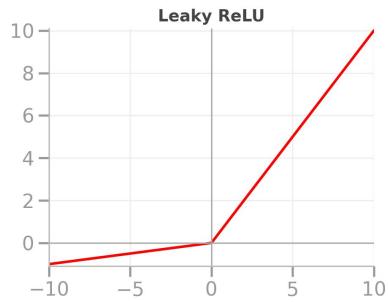
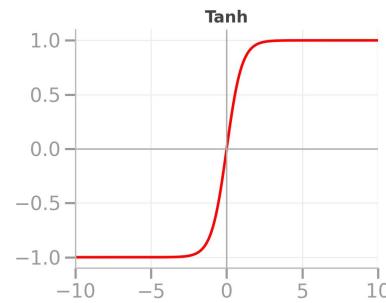
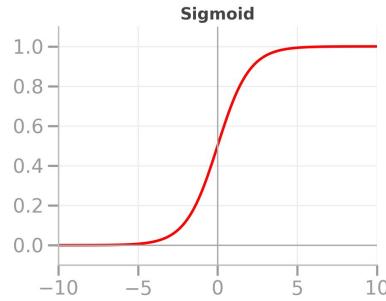
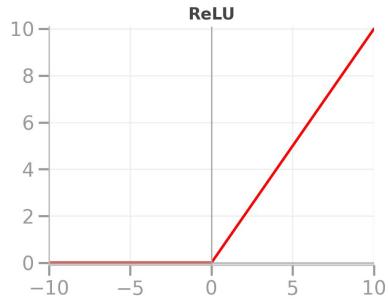


Non Linear Model

```
model = nn.Sequential(  
    nn.Linear(1, 5),  
    nn.ReLU(),  
    nn.Linear(5, 1)  
)
```



Activation Functions





<https://docs.pytorch.org/>



Specialization

Deep Learning Specialization

Build neural networks (CNNs, RNNs, LSTMs, Transformers) and apply them to speech recognition, NLP, and mor...

DeepLearning.AI

Specialization

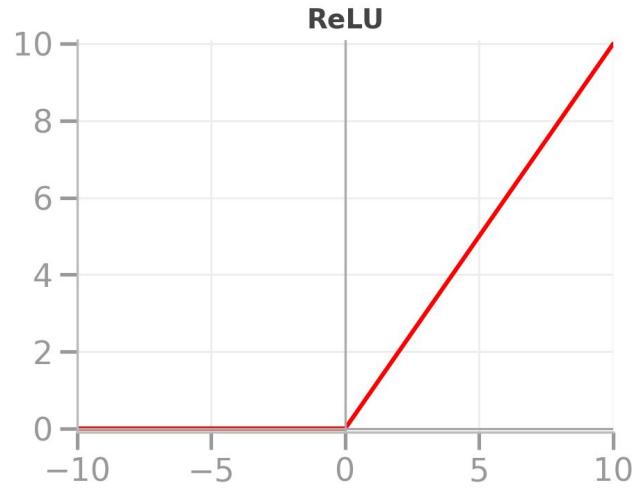
Machine Learning Specialization

Learn foundational AI concepts through an intuitive visual approach, then learn the code needed to...

DeepLearning.AI, Stanford Online

<https://wwwdeeplearning.ai/courses/deep-learning-specialization/>

<https://wwwdeeplearning.ai/courses/machine-learning-specialization/>





DeepLearning.AI

Tensors

Getting Started with PyTorch

You will learn how to...

- Read and understand tensor shapes
- Handle data types
- Create tensors
- Reshape and slice

Tensor Shapes

```
# Your delivery data
distances = torch.tensor([[3.0], [7.0], [12.0], [18.0], [22.0], [28.0]])
print(distances.shape)
```

Output:

```
C:>
torch.Size([6, 1])
```

Tensor Shapes

```
# Your delivery data
distances = torch.tensor([[3.0], [7.0], [12.0], [18.0], [22.0], [28.0]])
print(distances.shape)
```

Output:

```
C:>
torch.Size([6, 1])
```

batch size

Tensor Shapes

```
# Your delivery data
distances = torch.tensor([[3.0], [7.0], [12.0], [18.0], [22.0], [28.0]])
print(distances.shape)
```

Output:

```
C:>
torch.Size([6, 1])
```

features per sample

Tensor Shapes

```
# Your delivery data
distances = torch.tensor([[3.0], [7.0], [12.0], [18.0], [22.0], [28.0]])
print(distances.shape)
simple_model = nn.Linear(1, 1) # Expects 1 feature per sample
```

Output:

```
C:>
torch.Size([6, 1])
```

Tensor Shapes

```
# Your delivery data
distances = torch.tensor([[3.0], [7.0], [12.0], [18.0], [22.0], [28.0]])
print(distances.shape)
simple_model = nn.Linear(1, 1) # Expects 1 feature per sample

output = simple_model(distances) # Works!
```

Output:

```
C:>
torch.Size([6, 1])
```

Tensor Shapes

```
# Your delivery data
distances = torch.tensor([[3.0], [7.0], [12.0], [18.0], [22.0], [28.0]])
print(distances.shape)
simple_model = nn.Linear(1, 1) # Expects 1 feature per sample

output = simple_model(distances) # Works!
```

Output:

C:>
torch.Size([6, 1])

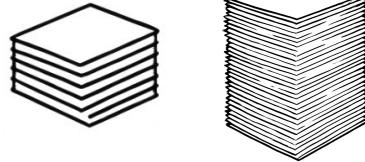
Tensor Shapes

```
# Your delivery data
distances = torch.tensor([[3.0], [7.0], [12.0], [18.0], [22.0], [28.0]])
print(distances.shape)
simple_model = nn.Linear(1, 1) # Expects 1 feature per sample

output = simple_model(distances)
```

C:>
torch.Size([6, 1])

batch size



Tensor Shapes

```
# Your delivery data
distances = torch.tensor([[3.0], [7.0], [12.0], [18.0], [22.0], [28.0]])
print(distances.shape)
simple_model = nn.Linear(1, 1) # Expects 1 feature per sample

output = simple_model(distances)
```

```
C:>
torch.Size([6, 1])
```

Tensor Shapes

```
# Your delivery data
distances = torch.tensor([[3.0, 7.0, 1.0], [18.0, 22.0, 2.0]])
print(distances.shape)
simple_model = nn.Linear(1, 1)

output = simple_model(distances)
```

distance hour weather

C:>
torch.Size([2, 3])

Tensor Shapes

```
# Your delivery data
distances = torch.tensor([[3.0, 7.0, 1.0], [18.0, 22.0, 2.0]])
print(distances.shape)
simple_model = nn.Linear(1, 1)

output = simple_model(distances)
```

C:>
torch.Size([2, 3]) RuntimeError

Tensor Shape Errors

RuntimeError: mat1 and mat2 shapes cannot be multiplied (2x3) and (1x1)

PyTorch will tell you what went wrong, but not always how to fix it.

Tensor Data Types

```
int_tensor = torch.tensor([1, 2, 3])  
print(int_tensor.dtype)
```

Output:

```
C:>  
torch.int64
```

Tensor Data Types

```
int_tensor = torch.tensor([1, 2, 3])  
print(int_tensor.dtype)
```

Output:

```
C:>  
torch.int64
```

```
float_tensor = torch.tensor([1.0, 2.0, 3.0])  
print(float_tensor.dtype)
```

Output:

```
C:>  
torch.float32
```

Tensor Data Types

```
float_tensor = torch.tensor([1, 2, 3], dtype=torch.float32)
```

Tensor Data Types

```
float_tensor = torch.tensor([1, 2, 3], dtype=torch.float32)  
float_tensor = int_tensor.float() # Convert to float32
```

Mixing Types

```
float_tensor = torch.tensor([1.0, 2.0, 3.0], dtype=torch.float32)

int_tensor = torch.tensor([1, 2, 3], dtype=torch.int64)

mixed_tensor = float_tensor + int_tensor

print(f"mixed type: {mixed_tensor.dtype}")

> mixed type: torch.float32
```

Tensor Data Types

- `float64` -> extra precision
- `int8` -> memory efficiency
- `float32` -> is often the default and works reliably

Creating Tensors

```
my_distances = [[3.0], [7.0], [12.0], [18.0]]  
tensor_from_list = torch.tensor(my_distances)
```

```
import numpy as np  
numpy_array = np.array([1.0, 2.0, 3.0])  
tensor_from_numpy = torch.from_numpy(numpy_array)
```

```
zeros = torch.zeros(3, 3)      # 3x3 tensor of zeros  
ones = torch.ones(2, 4)        # 2x4 tensor of ones  
random = torch.rand(5, 5)       # 5x5 tensor with random values
```

Reshaping Tensors

```
C:>  
torch.Size([6, 1])
```

batch size

Reshaping Tensors

```
# Your delivery data
single_distance = torch.tensor(25.0)
print(single_distance.shape)
simple_model = nn.Linear(1, 1)
```

C:>
torch.Size([]) ← *No dimensions!*

Model expects: torch.size([batch_size, 1])

Reshaping Tensors

```
# Your delivery data
single_distance = torch.tensor(25.0)

with_batch = single_distance.unsqueeze(0)
# torch.Size([1]) – added batch dim

ready_for_model = with_batch.unsqueeze(1)
# torch.Size([1, 1]) – now ready!
```

```
tensor([[25.0]])
```

Reshaping Tensors

```
# Your delivery data
single_distance = torch.tensor(25.0)

with_batch = single_distance.unsqueeze(0)
# torch.Size([1]) – added batch dim

ready_for_model = with_batch.unsqueeze(1)
# torch.Size([1, 1]) – now ready!
```

```
tensor([[25.0]])
```

```
# Your delivery data
squeezed = ready_for_model.squeeze()
```

```
tensor(25.0)
```

```
print(my_tensor.shape)
```

Indexing and Slicing

```
# Your predictions from the delivery model
predictions = torch.tensor([[14.9], [24.1], [35.6], [45.2]])
```

Indexing and Slicing

```
# Your predictions from the delivery model
predictions = torch.tensor([[14.9], [24.1], [35.6], [45.2]])

# Get the first prediction    Output:
first = predictions[0]
print(first)
```

C:>
tensor([14.9])

Indexing and Slicing

```
# Your predictions from the delivery model
predictions = torch.tensor([[14.9], [24.1], [35.6], [45.2]])

# Get the first prediction
first = predictions[0]
print(first)

# Get first three predictions  Output:
first_three = predictions[:3]
print(first_three)
```

C:>
tensor([[14.9000],
 [24.1000],
 [35.6000]])

Indexing and Slicing

```
# Your predictions from the delivery model
predictions = torch.tensor([[14.9], [24.1], [35.6], [45.2]])

# Get the first prediction
first = predictions[0]
print(first)

# Get first three predictions
first_three = predictions[:3]
print(first_three)

# Use .item() to get a Python number
value = predictions[0].item()
```

Indexing and Slicing

```
# Your predictions from the delivery model
predictions = torch.tensor([[14.9], [24.1], [35.6], [45.2]])

# Get the first prediction
first = predictions[0]
print(first)

# Get first three predictions
first_three = predictions[:3]
print(first_three)

# Use .item() to get a Python number
value = predictions[0].item()
```

```
C:>
14.899999618530273
<class 'float'>
```

Indexing and Slicing

```
# Multiple features per sample
data = torch.tensor([[3.0, 8.0, 1.0],      # distance, hour, weather
                     [7.0, 17.0, 2.0],
                     [12.0, 12.0, 1.0]])
```

```
distances = data[:, 0]
```

Output:

```
C:>
tensor([ 3.,  7., 12.])
```

Tensors

- Shapes
- Types
- Reshaping
- Indexing



DeepLearning.AI

Tensor Math and Broadcasting

Getting Started with PyTorch

Tensor Math

```
distances = [[3.0], [7.0], [12.0]]  
weight = 2.3  
bias = 8.0  
  
weight * distances + bias  
  
[[3.0], [7.0], [12.0]]
```

Tensor Math

```
distances = [[3.0], [7.0], [12.0]]
```

```
weight = 2.3
```

```
bias = 8.0
```

```
weight * distances + bias
```

```
[[2.3*3.0], [2.3*7.0], [2.3*12.0]]
```

Tensor Math

```
distances = [[3.0], [7.0], [12.0]]
```

```
weight = 2.3
```

```
bias = 8.0
```

```
weight * distances + bias
```

```
[[2.3*3.0 + 8.0], [2.3*7.0 + 8.0], [2.3*12.0 + 8.0]]
```

Tensor Math

Scalars

$$2.3 * [[3.0], [7.0], [12.0]] = [[2.3*3.0], [2.3*7.0], [2.3*12.0]]$$

Tensors with same Shape

$$[[1.0], [2.0], [3.0]] * [[3.0], [7.0], [12.0]] = [[1.0*3.0], [2.0*7.0], [3.0*12.0]]$$

Tensor Math

```
torch.tensor([[3.0,     8.0,     1.0],    ← 1
              [7.0,     17.0,    2.0],   ← 2
              [12.0,    12.0,    1.0]])) ← 3
```

Tensor Math

distance hour weather

```
torch.tensor([[3.0,            8.0,            1.0],
              [7.0,            17.0,          2.0],
              [12.0,          12.0,          1.0]]))
```

Tensor Math

```
*1.1      *1.0      *5.0  
↓          ↓          ↓  
torch.tensor([[3.0, 8.0, 1.0],  
             [7.0, 17.0, 2.0],  
             [12.0, 12.0, 1.0]]))
```

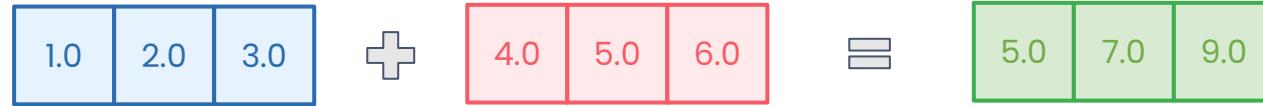
Tensor Math

```
*1.1      *1.0      *5.0  
↓         ↓         ↓  
torch.tensor([[3.0,     8.0,     1.0],           [[1.1,     1.0,     5.0],  
          [7.0,     17.0,    2.0],    *   [1.1,     1.0,     5.0],  
          [12.0,    12.0,    1.0]]))           [1.1,     1.0,     5.0]])
```

Tensor Math

```
*1.1      *1.0      *5.0  
↓         ↓         ↓  
torch.tensor([[3.0,     8.0,     1.0],  
              [7.0,     17.0,    2.0],    *  [[1.1,     1.0,     5.0]]  
              [12.0,    12.0,    1.0]]))
```

Broadcasting



`[[1.0, 2.0, 3.0]]`

`shape = (1, 3)`

`[[4.0, 5.0, 6.0]]`

`shape = (1, 3)`

`[[5.0, 7.0, 9.0]]`

`shape = (1, 3)`

Broadcasting

$$\begin{bmatrix} 1.0 & 2.0 & 3.0 \end{bmatrix} + 5.0 = \begin{bmatrix} 6.0 & 7.0 & 8.0 \end{bmatrix}$$

`[[1.0, 2.0, 3.0]]`

`shape = (1, 3)`

5.0

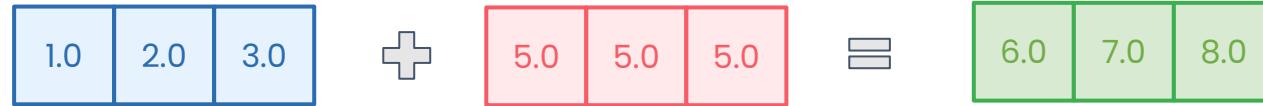
`5.0`

`shape = ()`

`[[6.0, 7.0, 8.0]]`

`shape = (1, 3)`

Broadcasting



`[[1.0, 2.0, 3.0]]`

`shape = (1, 3)`

`[[5.0, 5.0, 5.0]]`

`shape = (1, 3)`

`[[6.0, 7.0, 8.0]]`

`shape = (1, 3)`

Broadcasting

1.0	2.0	3.0
-----	-----	-----



5.0



`[[1.0, 2.0, 3.0]]`

`shape = (1, 3)`

`[[5.0]]`

`shape = (1,1)`

Mismatched Shapes

Broadcasting

1.0	2.0	3.0
-----	-----	-----



5.0



?

`[[1.0, 2.0, 3.0]]`

`shape = (1, 3)`

`[[5.0]]`

`shape = (1, 1)`

Broadcasting

1.0	2.0	3.0
-----	-----	-----



5.0



`[[1.0, 2.0, 3.0]]`

`shape = (1, 3)`

`[[5.0]]`

`shape = (1, 3)`

Broadcasting

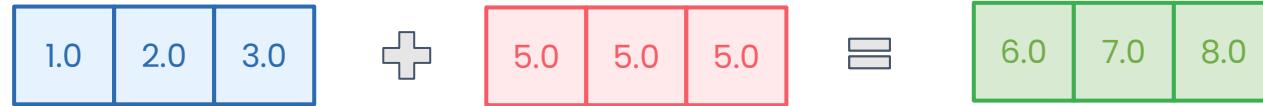
$$\begin{array}{|c|c|c|} \hline 1.0 & 2.0 & 3.0 \\ \hline \end{array} + \begin{array}{|c|c|c|} \hline 5.0 & 5.0 & 5.0 \\ \hline \end{array} = ?$$

`[[1.0, 2.0, 3.0]]` `[[5.0, 5.0, 5.0]]`

`shape = (1, 3)`

`shape = (1, 3)`

Broadcasting



`[[1.0, 2.0, 3.0]]`

`shape = (1, 3)`

`[[5.0, 5.0, 5.0]]`

`shape = (1, 3)`

`[[6.0, 7.0, 8.0]]`

`shape = (1, 3)`

Broadcasting

$$\begin{array}{c} \begin{array}{|c|c|c|} \hline 1.0 & 2.0 & 3.0 \\ \hline \end{array} + \begin{array}{|c|} \hline 4.0 \\ \hline 5.0 \\ \hline 6.0 \\ \hline \end{array} = ? \end{array}$$

`[[1.0, 2.0, 3.0]]`

shape = `(1, 3)`

`[[4.0], [5.0], [6.0]]`

shape = `(3, 1)`

Broadcasting

$$\begin{array}{c} \begin{array}{|c|c|c|} \hline 1.0 & 2.0 & 3.0 \\ \hline \end{array} + \begin{array}{|c|} \hline 4.0 \\ \hline 5.0 \\ \hline 6.0 \\ \hline \end{array} = ? \end{array}$$

`[[1.0, 2.0, 3.0]]`

shape = (1, 3)

`[[4.0], [5.0], [6.0]]`

shape = (3, 1)

Broadcasting

1.0	2.0	3.0
1.0	2.0	3.0
1.0	2.0	3.0



4.0
5.0
6.0



```
[[1.0, 2.0, 3.0],  
 [1.0, 2.0, 3.0],  
 [1.0, 2.0, 3.0]]
```

shape = (3,3)

```
[[4.0], [5.0], [6.0]]
```

shape = (3,1)

Broadcasting

1.0	2.0	3.0
1.0	2.0	3.0
1.0	2.0	3.0



4.0
5.0
6.0



```
[[1.0, 2.0, 3.0],  
 [1.0, 2.0, 3.0],  
 [1.0, 2.0, 3.0]]
```

shape = (3 , 3)

```
[[4.0], [5.0], [6.0]]
```

shape = (3 , 1)

Broadcasting

1.0	2.0	3.0
1.0	2.0	3.0
1.0	2.0	3.0



4.0	4.0	4.0
5.0	5.0	5.0
6.0	6.0	6.0



?

```
[[1.0, 2.0, 3.0],  
 [1.0, 2.0, 3.0],  
 [1.0, 2.0, 3.0]]
```

shape = (3 , 3)

```
[[4.0, 4.0, 4.0],  
 [5.0, 5.0, 5.0],  
 [6.0, 6.0, 6.0]]
```

shape = (3 , 3)

Broadcasting

1.0	2.0	3.0
1.0	2.0	3.0
1.0	2.0	3.0



4.0	4.0	4.0
5.0	5.0	5.0
6.0	6.0	6.0



5.0	6.0	7.0
6.0	7.0	8.0
7.0	8.0	9.0

```
[[1.0, 2.0, 3.0],  
 [1.0, 2.0, 3.0],  
 [1.0, 2.0, 3.0]]
```

shape = (3 , 3)

```
[[4.0, 4.0, 4.0],  
 [5.0, 5.0, 5.0],  
 [6.0, 6.0, 6.0]]
```

shape = (3 , 3)

```
[[5.0, 6.0, 7.0],  
 [6.0, 7.0, 8.0],  
 [7.0, 8.0, 9.0]]
```

shape = (3 , 3)

Tensor Math

```
torch.tensor([[3.0,     8.0,     1.0],           [[1.1,     1.0,     5.0],  
     [7.0,     17.0,    2.0],    *   [1.1,     1.0,     5.0],  
     [12.0,    12.0,    1.0]]))      [1.1,     1.0,     5.0]])
```

Tensor Math

```
torch.tensor([[3.0,     8.0,     1.0],  
             [7.0,     17.0,    2.0],    *    [[1.1,     1.0,     5.0]]  
             [12.0,    12.0,    1.0]]))
```